
Parallel Programming

Course Introduction





Professor Yi-Ping You (游逸平)

Department of Computer Science

<http://www.cs.nctu.edu.tw/~ypyou/>



Parallel Processing/Computing



Finding the integer solutions for $x^3+y^3+z^3=3$

- The first two solutions might be obvious
$$1^3 + 1^3 + 1^3 = 3$$
$$4^3 + 4^3 + (-5)^3 = 3$$
- Is it even possible to know whether other solutions for 3 exist? (Louis Mordell, 1953)



- ✚ Roughly 4 million tasks (three hours/task)
- ✚ Using Charity Engine's grid comprised over 400,000 computers around the world

<https://phys.org/news/2021-03-sum-cubes-puzzle-solution.html>

<https://www.pnas.org/doi/10.1073/pnas.2022377118>

<https://github.com/AndrewVSutherland/SumsOfThreeCubes>



Motivation



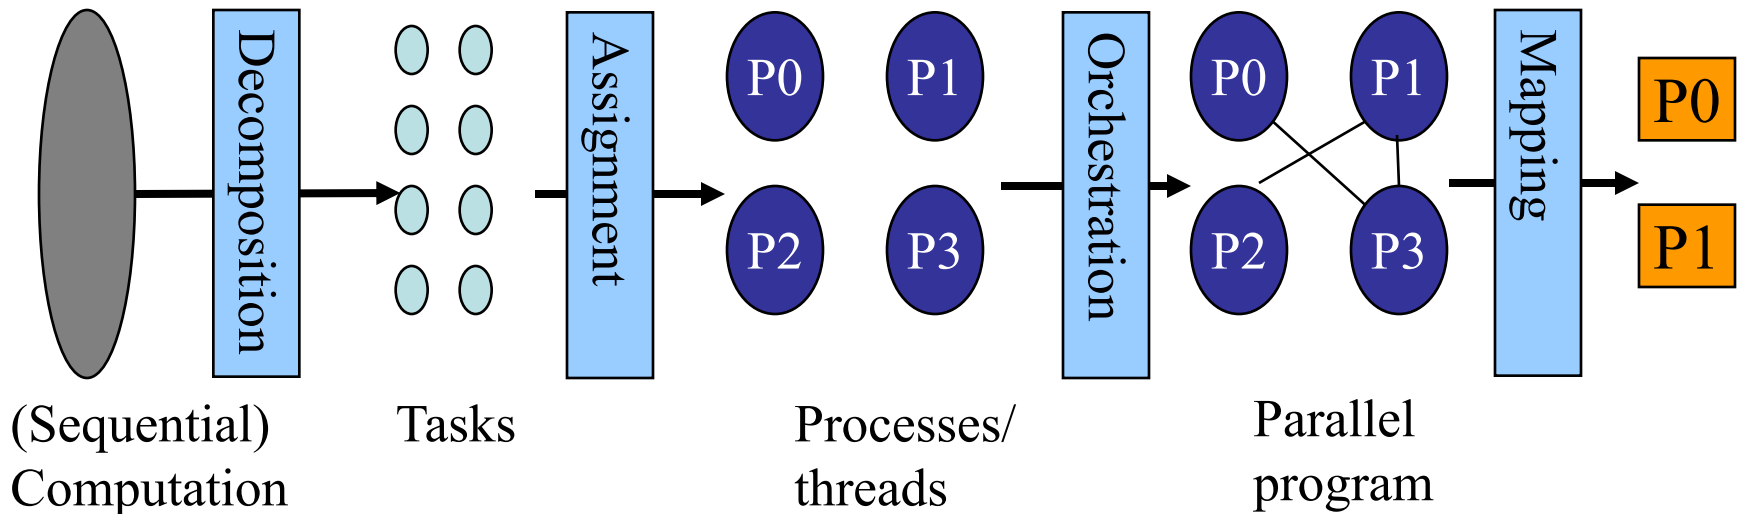
Creating a Parallel Program

- In theory, can be done by programmer, compiler, run-time system, or OS
- In practice, parallel programs are created with
 - ⊕ Explicitly parallel language (e.g., High Performance Fortran)
 - ⊕ Library for implementing a programming model
 - ◆ Shared-memory library (Pthreads, OpenMP)
 - ◆ Distributed-memory library (Message Passing Interface)
 - ◆ Heterogeneous-programming library (CUDA, OpenCL)
 - ◆ Cluster-based library (MapReduce)




Steps for Creating a Parallel Program

- **Decomposition** into tasks
- **Assignment** of tasks to processes/threads
- **Orchestration** of data access, communication, etc.
- **Mapping** processes to processors




Shared-Memory Systems

- Multiple processors can operate independently but share the same memory resources




Distributed-Memory Systems

- Processors have their own local memory
- Memory addresses in one processor do not map to another processor
 - ⊕ So there is no concept of global address space across all processors



Hybrid Distributed-Shared Systems

- The distributed memory component is the networking of multiple shared memory machines, which know only about their own memory - not the memory on another machine



Administrative Stuff

- Course information

- Parallel Programing
- Credit: 3
- Time: Thursdays 9:00-12:00
- Place: EDB27

- Course website

- <https://pp-f23.github.io/>
- The URL is also provided on my Web page
 - ◆ Authorization required to access lecture slides



Prerequisites

- This course assumes that you
 - ✦ write good code in C/C++ and
 - ✦ are familiar with the Linux environment
- Requirements
 - ✦ Have taken a C/C++ course or have a fair amount of practical experience with C/C++ programming
 - ✦ Know how to connect to remote machines with `ssh`, `scp`, etc.
 - ✦ Know how to work on Linux systems (i.e., basic commands such as `cd`, `rm`, and running executables and GNU tools such as `make` and `gcc`)



Aims of This Course

- The skills and knowledge needed to develop applications using parallel programming models



Lecture Topics


- Background
 - Parallel and distributed programming
 - Introduction to parallel hardware and software
- Shared-memory programming
 - Pthreads and OpenMP
- Distributed-memory programming
 - MPI and MapReduce
- GPGPU programming
 - CUDA and OpenCL



MPICH





Open MPI



Textbook

- Peter Pacheco, An Introduction to Parallel Programming, Morgan Kaufmann; 1 edition (January 21, 2011), 2nd edition (March 15, 2020)



Grading

- Grades will be assigned based on
 - ⊕ Homework assignments (70%)
 - ◆ 6 assignments related to parallel programming
 - ◆ **Slackers beware!**
 - The penalty for late homework is **15% per day** (weekends count as 1 day).
 - ◆ **NO PLAGIARISM!**
 - Homework assignments must be individual work
 - ⊕ Course project (30%)
 - ◆ 3 students form a group to work on development of parallel applications
 - ◆ Proposal (4%)
 - ◆ Final oral presentation (13%)
 - ◆ Final report (13%)
- These weights are subject to minor variation



Project Schedule

- Group registration due on October 5, 2023 (Week 4)

■ <https://ppt.cc/fKxpqx>



- Project proposal due on November 2, 2023 (Week 8)
- Presentation slides due by 23:59 day prior to your presentation
- Final report and source codes due on January 8, 2024 (Week 18)



HW0: A Warming-Up Assignment

- To assess whether you are familiar with Makefile and C/C++ programming
- See the URL below for details
 - ✦ <https://PP-f23.github.io/HW0/>



Discussion Forum

- The E3 system is used for discussion

