

Parallel Programming

Shared-Memory Programming with OpenMP (Part I)

Professor Yi-Ping You (游逸平)
Department of Computer Science
<http://www.cs.nctu.edu.tw/~ypyou/>



Outline

- Introduction to OpenMP
- Compiler Directives
 - ❖ Creating Threads
 - ❖ Synchronization
 - ❖ Worksharing Constructs
 - ❖ More on Synchronization
 - ❖ Variable Scopes (shared/private)
- Run-time Libraries and Environment Variables



Motivation

- Thread libraries are hard to use
 - Pthreads threads have many library calls for initialization, synchronization, thread creation, condition variables, etc.
 - Programmer must code with multiple threads in mind
- Synchronization between threads introduces a new dimension of program correctness
- Wouldn't it be nice to write serial programs and somehow parallelize them “automatically”?
 - OpenMP can parallelize many serial programs with relatively few annotations that specify parallelism and independence
 - It is not automatic: you can still make errors in your annotations



Introduction to OpenMP

■ What is OpenMP?

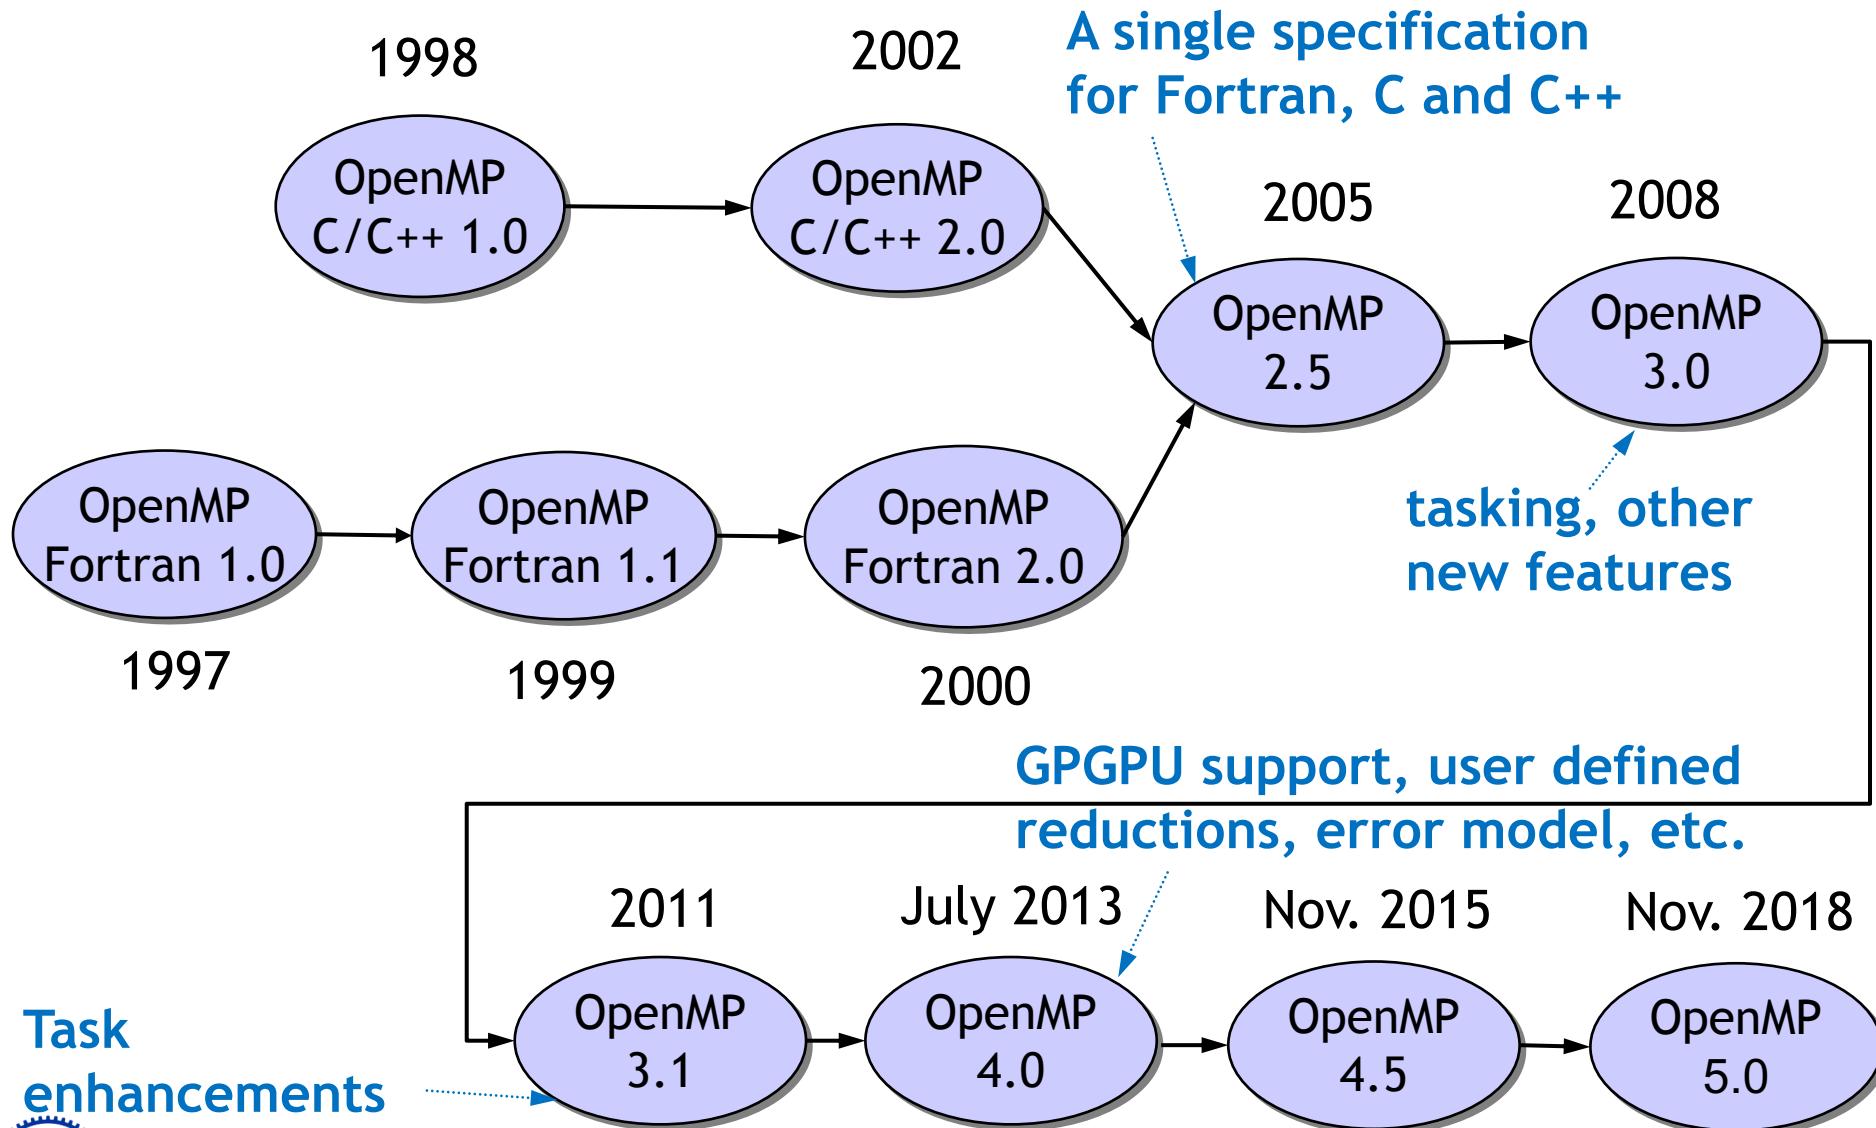
- ❖ Open specification for Multi-Processing
- ❖ “Standard” API for defining multi-threaded shared-memory programs
 - ◆ Standardizes loop-level & task parallelism
 - ◆ Standardizes ~ 20 years of compiler-directed threading experience

■ High-level API

- ❖ Preprocessor (compiler) directives (~ 80%)
- ❖ Library Calls (~ 19%)
- ❖ Environment Variables (~ 1%)



OpenMP Release History



A Programmer's View of OpenMP

- OpenMP is a portable, threaded, shared-memory programming specification with “light” syntax
 - ⊕ Exact behavior depends on OpenMP implementation!
 - ⊕ Requires compiler support (C/C++ or Fortran)
- OpenMP will:
 - ⊕ Allow programmers to incrementally parallelize existing serial programs
 - ⊕ Allow a programmer to separate a program into serial regions and parallel regions, rather than T concurrently-executing threads
 - ⊕ Hide stack management
 - ⊕ Provide synchronization constructs
- OpenMP will NOT:
 - ⊕ Parallelize automatically
 - ⊕ Guarantee speedup
 - ⊕ Provide freedom from data races

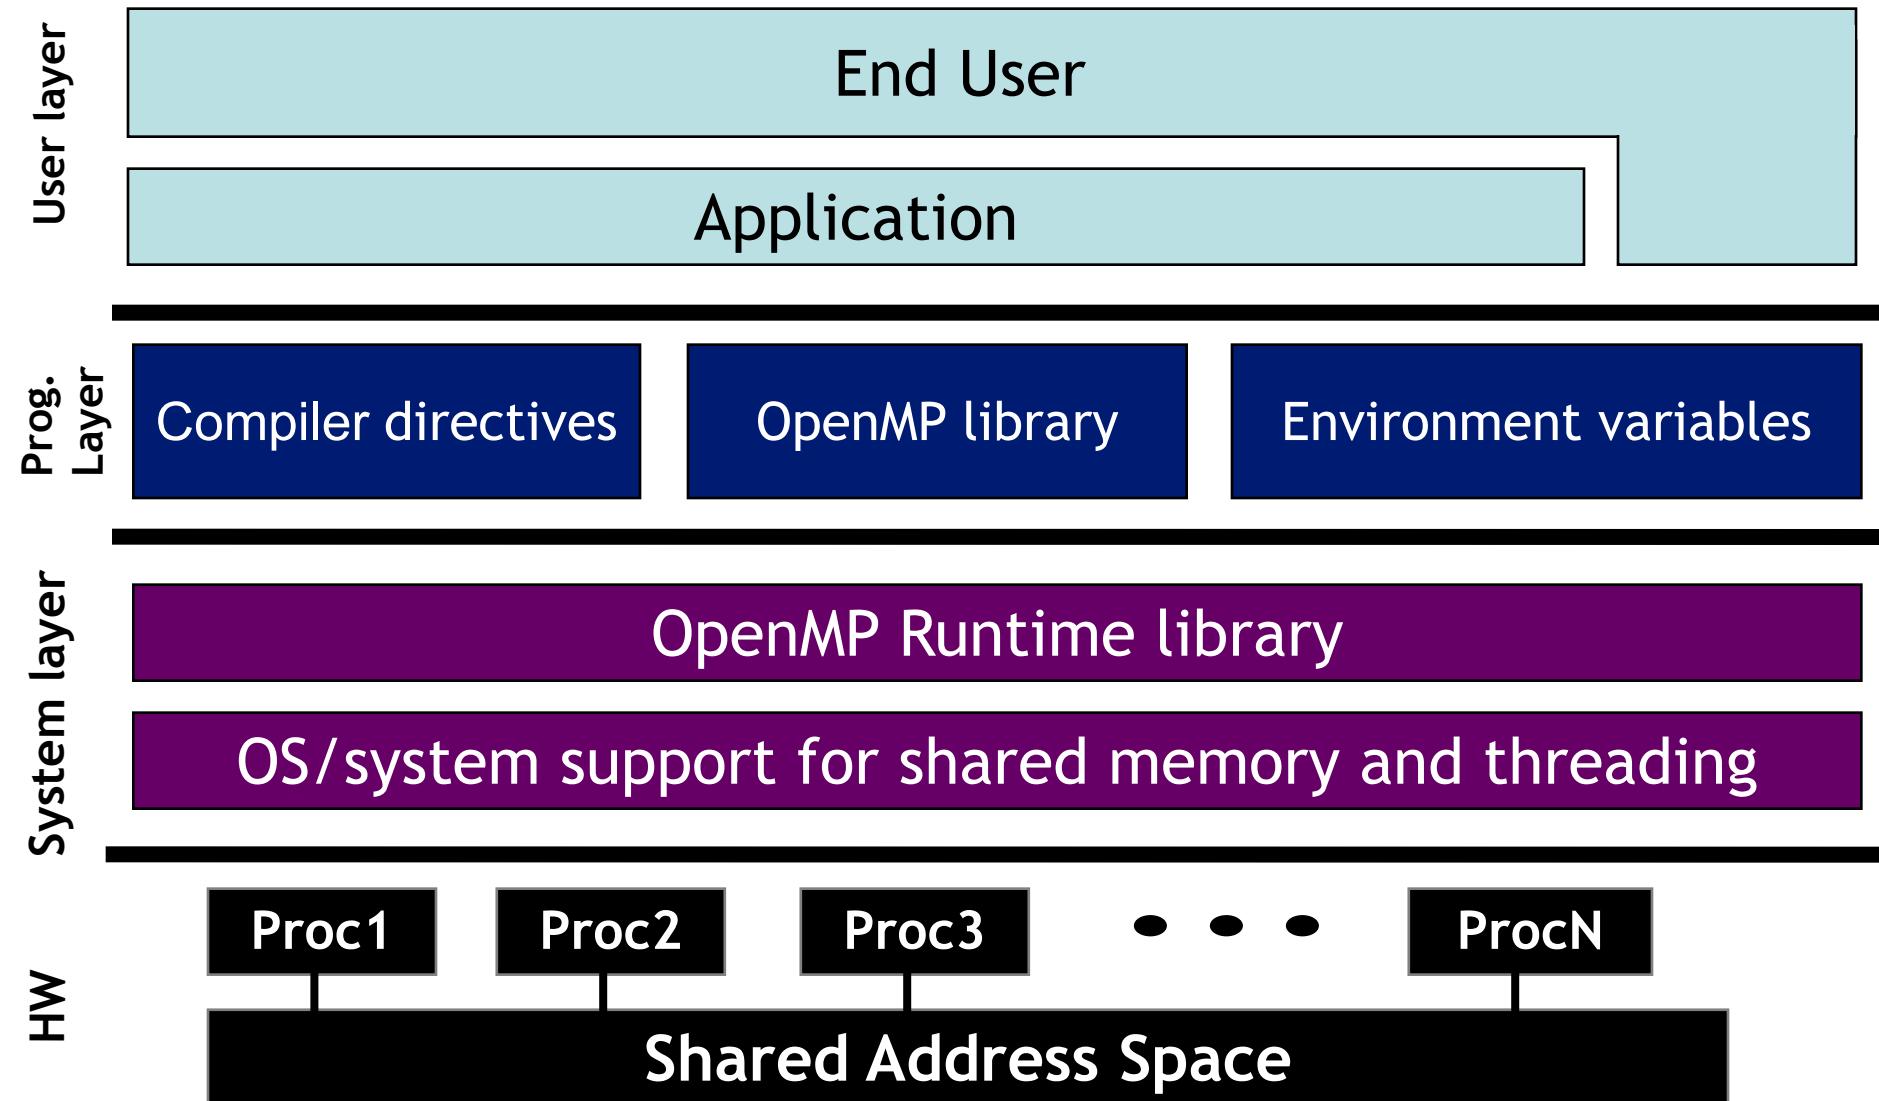


What is OpenMP?

- Three components:
 1. Set of **compiler directives** (80%) for
 - ❖ creating teams of threads
 - ❖ sharing the work among threads
 - ❖ synchronizing the threads
 2. **Library routines** (19%) for setting and querying thread attributes
 3. **Environment variables** (1%) for controlling run-time behavior of the parallel program



OpenMP Basic Defs: Solution Stack



OpenMP Core Syntax

- Most of the constructs in OpenMP are compiler directives
 - ◆ **#pragma omp construct [clause [clause]...]**
 - ◆ E.g., #pragma omp parallel num_threads(4)
- Function prototypes and types in the file:
 - ◆ `#include <omp.h>`
- Most OpenMP constructs apply to a “structured block”
 - ◆ Structured block: a block of one or more statements with one point of entry at the top and one point of exit at the bottom (**single-entry, single-exit**)
 - ◆ It’s OK to have an `exit()` within the structured block



OpenMP Example: Hello World

- Write a multithreaded program where each thread prints “hello world”

```
int thread_count = 3;  
void main() {  
  
    printf("Hello from thread %d of %d",  
           ID, thread_count);  
  
}
```



OpenMP Example: Hello World

- Tell the compiler to pack code into a function, fork the threads, and join when done ...

```
#include "omp.h"           OpenMP header file
int thread_count = 3;
void main() {
#pragma omp parallel      Parallel region with default
{                         number of threads
    int ID = omp_get_thread_num();  Runtime library function to
    printf("Hello from thread %d of %d",
           ID, thread_count);
}
}                           End of the parallel region
```

Sample Output:
Hello from thread 0 of 3
Hello from thread 1 of 3
Hello from thread 2 of 3



Compilation

- To compile with gcc, we need to include the `-fopenmp` option
 - ✿ `$ gcc -fopenmp -o omp_hello
omp_hello.c`
- Options vary among different compilers

Platform	Compile option
Linux and OS X	<code>gcc -fopenmp</code>
PGI Linux	<code>pgcc -mp</code>
Intel windows	<code>icl /Qopenmp</code>
Intel Linux and OS X	<code>icpc -openmp</code>



OpenMP Overview: How do threads interact?

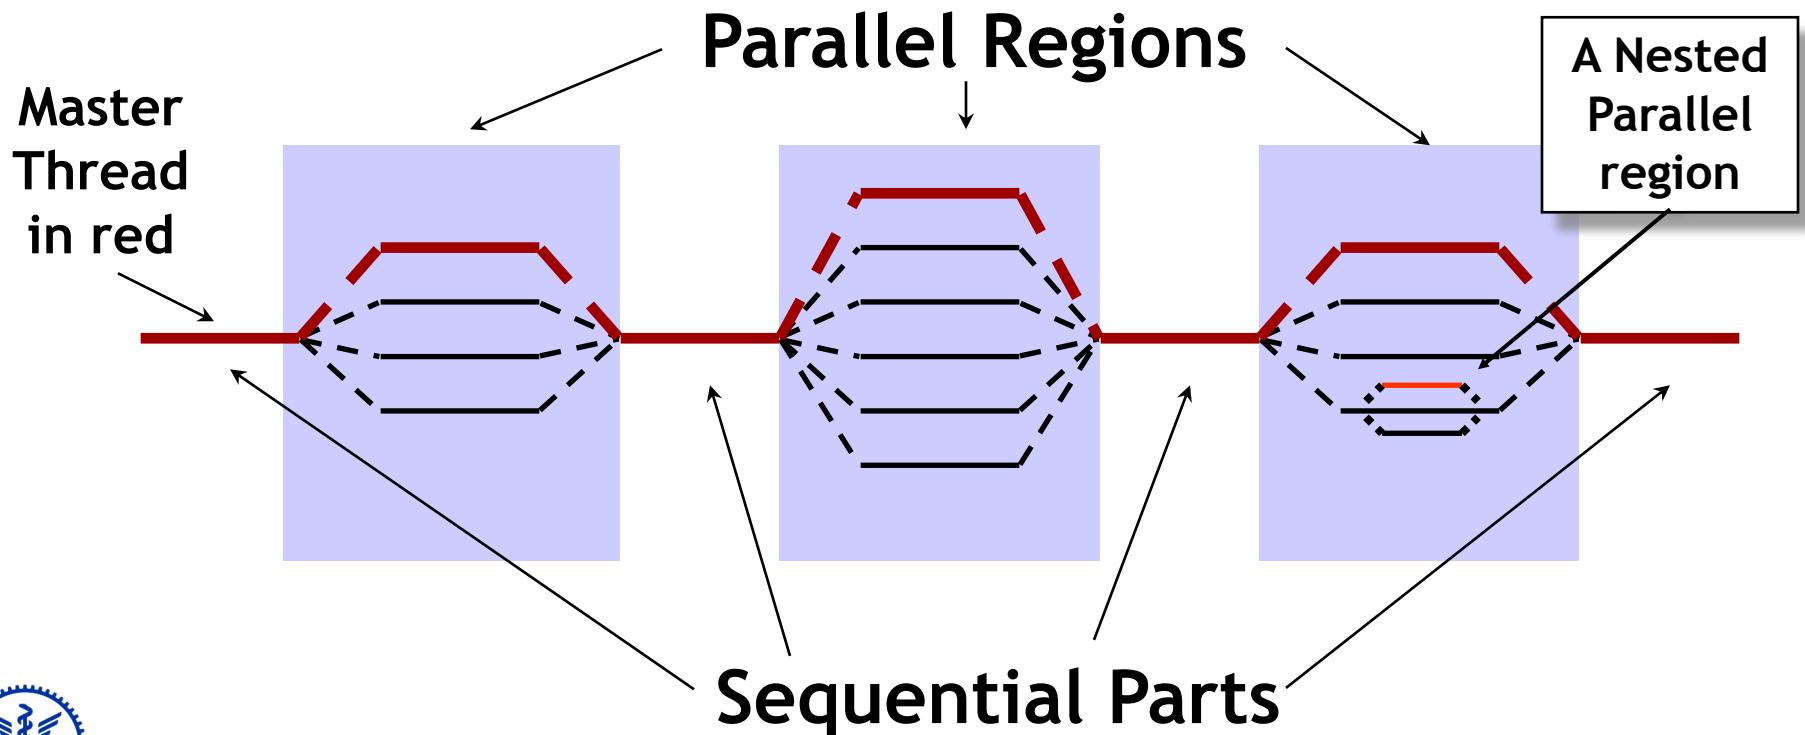
- OpenMP is a multi-threading, shared address model
 - Threads communicate by sharing variables
- Unintended sharing of data causes race conditions
 - race condition: when the program's outcome changes as the threads are scheduled differently
- To control race conditions
 - Use synchronization to protect data conflicts
- Synchronization is expensive so
 - Change how data is accessed to minimize the need for synchronization



OpenMP Execution Model

Fork-Join Parallelism:

- Master thread spawns a team of threads as needed
- Parallelism added incrementally until performance are met: i.e. the sequential program evolves into a parallel program



Outline

- Introduction to OpenMP
- Compiler Directives
 - ❖ Creating Threads
 - ❖ Synchronization
 - ❖ Worksharing Constructs
 - ❖ More on Synchronization
 - ❖ Variable Scopes (shared/private)
- Run-time Libraries and Environment Variables



Thread Creation: Parallel Regions (1/2)

- You create threads in OpenMP with the parallel construct
- For example, to create a 4-thread parallel region:

```
double A[1000];
omp_set_num_threads(4);
#pragma omp parallel
{
    int ID = omp_get_thread_num();
    pooh(ID, A);
}
```

Runtime function to request
a certain number of threads

Runtime library function to
return a thread ID

Each thread executes a copy of the code within the structured block

- Each thread calls `pooh (ID, A)` for $ID = 0$ to 3



Thread Creation: Parallel Regions (2/2)

- You create threads in OpenMP with the parallel construct
- For example, to create a 4-thread parallel region:

```
double A[1000];
```

Clause to request a certain number of threads

```
#pragma omp parallel num_threads(4)
```

```
{
```

```
    int ID = omp_get_thread_num();
```

```
    pooh(ID, A);
```

```
}
```

Runtime library function to return a thread ID

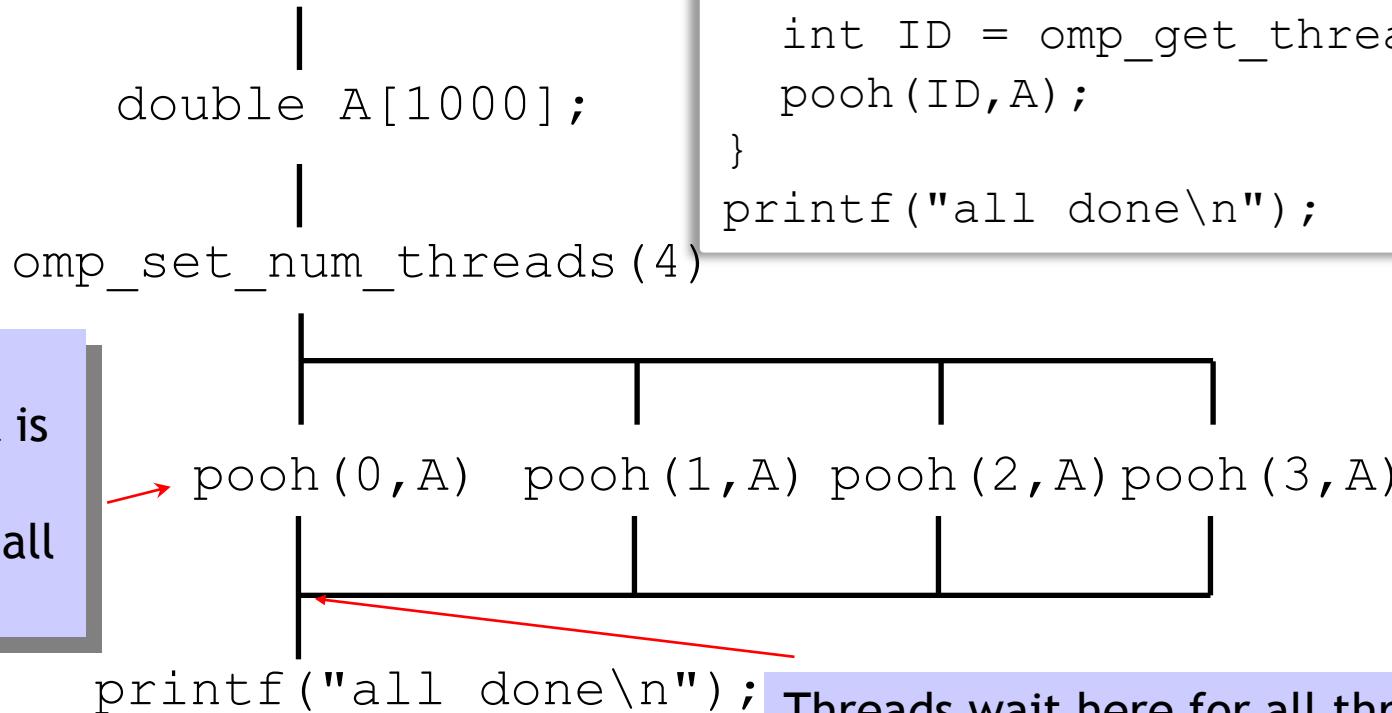
Each thread executes a copy of the code within the structured block

- Each thread calls `pooh (ID, A)` for $ID = 0$ to 3



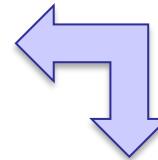
Thread Creation: Parallel Regions Example

- Each thread executes the same code redundantly



OpenMP: What the Compiler Does?

```
#pragma omp parallel num_threads(4)
{
    foobar();
}
```



- The OpenMP compiler generates code logically analogous to that on the right of this slide, given an OpenMP pragma such as that on the top-left
- All known OpenMP implementations use a thread pool so full cost of threads creation and destruction is not incurred for each parallel region
- Only three threads are created because the last parallel section will be invoked from the parent thread

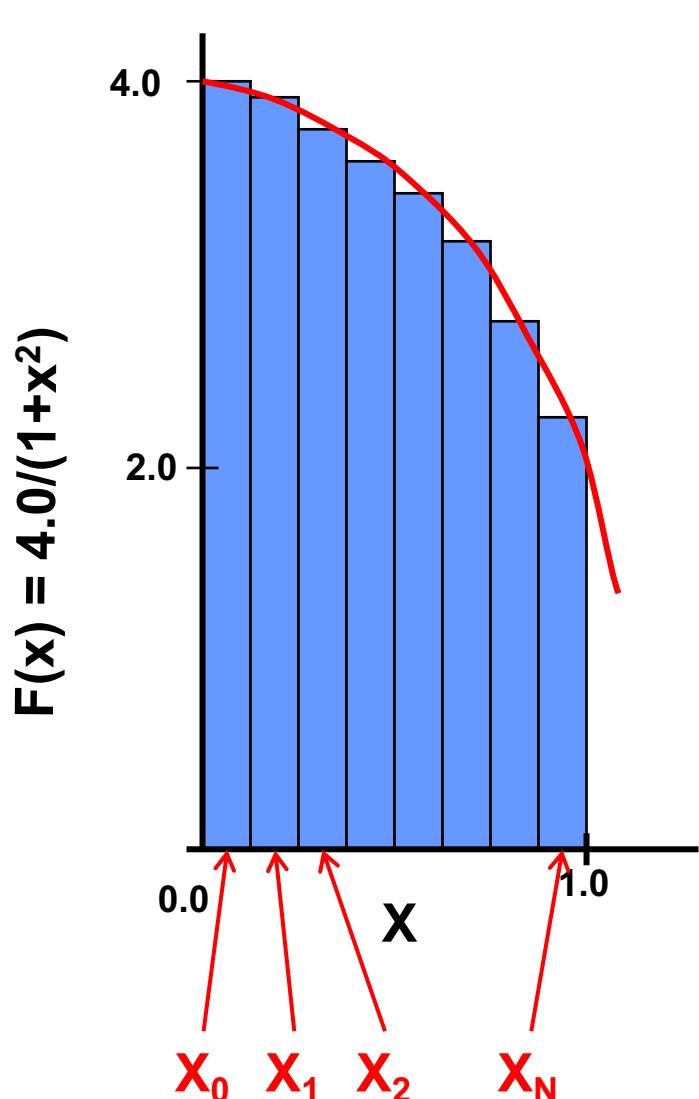
```
void thunk ()
{
    foobar();
}

pthread_t tid[4];
for (int i = 1; i < 4; ++i)
    pthread_create (
        &tid[i], 0, thunk, 0);
thunk();

for (int i = 1; i < 4; ++i)
    pthread_join(tid[i],
                 NULL);
```



Estimation of Pi: Numerical Integration



- Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

- We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i



Serial Pi Program

```
static long num_steps = 100000;
double step;
void main ()
{
    int i;
    double x, pi, sum = 0.0;

    step = 1.0/(double) num_steps;
    x = 0.5*step;
    for (i=0; i< num_steps; i++) {
        sum = sum + 4.0/(1.0+x*x);
        x += step;
    }
    pi = step * sum;
}
```



Pi Program: Transform into a Pthreads Program

```
static long num_steps = 100000;  
double step;  
void main ()  
{  
    int i;  
    double x, pi, sum = 0.0;
```

Variable to accumulate thread results
must be shared

Package
this into a
function

```
step = 1.0 / (double) num_steps;  
  
for (i=0; i< num_steps; i++) {  
    sum = sum + 4.0 / (1.0+x*x);  
    x = (i+0.5)*step;  
}  
pi = step * sum;
```

Assign loop
iterations to threads

Remove loop-carried
dependence

Assure safe update to pi



Numerical Integration: Pthreads (1/2)

```
#define NUMSTEPS 10000000
#define NUMTHREADS 4
double step = 0.0, Pi = 0.0;
pthread_mutex_t gLock;
int main() {
    pthread_t thrds[NUMTHREADS];
    int tRank[NUMTHREADS], i;
    pthread_mutex_init(&gLock, NULL);
    step = 1.0 / NUMSTEPS;
    for ( i = 0; i < NUMTHREADS; ++i ) {
        tRank[i] = i;
        pthread_create(&thrds[i], NULL, Func, (void) &tRank[i]);
    }
    for ( i = 0; i < NUMTHREADS; ++i )
        pthread_join(thrds[i], NULL);
    pthread_mutex_destroy(&gLock);
    printf("Computed value of Pi: %12.9f\n", Pi );
    return 0;
}
```

Global variables (in data section)

Initialize the mutex variable

Create (fork) the threads

Post a join for each thread



Numerical Integration: Pthreads (2/2)

```
void *Func(void *pArg) {  
    int myRank = *((int *)pArg);  
    double partialSum = 0.0, x;  
    for (int i = myRank; i < NUMSTEPS; i += NUMTHREADS) {  
        x = (i + 0.5f) * step;  
        partialSum += 4.0f / (1.0f + x*x);  
    }  
    pthread_mutex_lock(&gLock);  
    Pi += partialSum * Step;  
    pthread_mutex_unlock(&gLock);  
  
    return 0;  
}
```

Local variables (on each thread's own stack). partialSum to avoid competition

Use the global mutex to ensure exclusive access to Pi



A Simple Parallel Pi Program using OpenMP

```
#include <omp.h>
static long num_steps = 100000; double step;
#define NUM_THREADS 2
void main () {
    int i, nthreads;
    double x, pi, sum[NUM_THREADS];
    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int i, id, nthrds;
        double x;
        id = omp_get_thread_num();
        nthrds= omp_get_num_threads();
        if (id == 0) nthreads = nthrds;
        for (i=id, sum[id]=0.0; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            sum[id] += 4.0/(1.0+x*x);
        }
    }
    for(i=0, pi=0.0; i < nthreads; i++)
        pi += sum[i] * step;
}
```



SPMD: Single Program Multiple Data

- Run the same program on P processing elements where P can be arbitrarily large
- Use the rank—an ID ranging from 0 to $(P-1)$ —to select between a set of tasks and to manage any shared data structures

This pattern is very general and has been used to support most (if not all) the algorithm strategy patterns

It is probably the most commonly used pattern in the history of parallel programming



Results*

- Original Serial pi program with 100,000,000 steps ran in 1.83 seconds

```
#include <omp.h>
static long num_steps = 100000; double step;
#define NUM_THREADS 2
void main () {
    int i, nthreads;
    double x, pi, sum[NUM_THREADS];
    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
    int i, id, nthrds;
    double x;
    id = omp_get_thread_num();
    nthrds= omp_get_num_threads();
    if (id == 0) nthreads = nthrds;
    for (i=id, sum[id]=0.0; i< num_steps; i=i+nthrds) {
        x = (i+0.5)*step;
        sum[id] += 4.0/(1.0+x*x);
    }
}
for(i=0, pi=0.0; i < nthreads; i++)
    pi += sum[i] * step;
}
```

Threads	1 st SPMD	Eff.
1	1.86	0.98
2	1.03	0.89
3	1.08	0.56
4	0.97	0.47

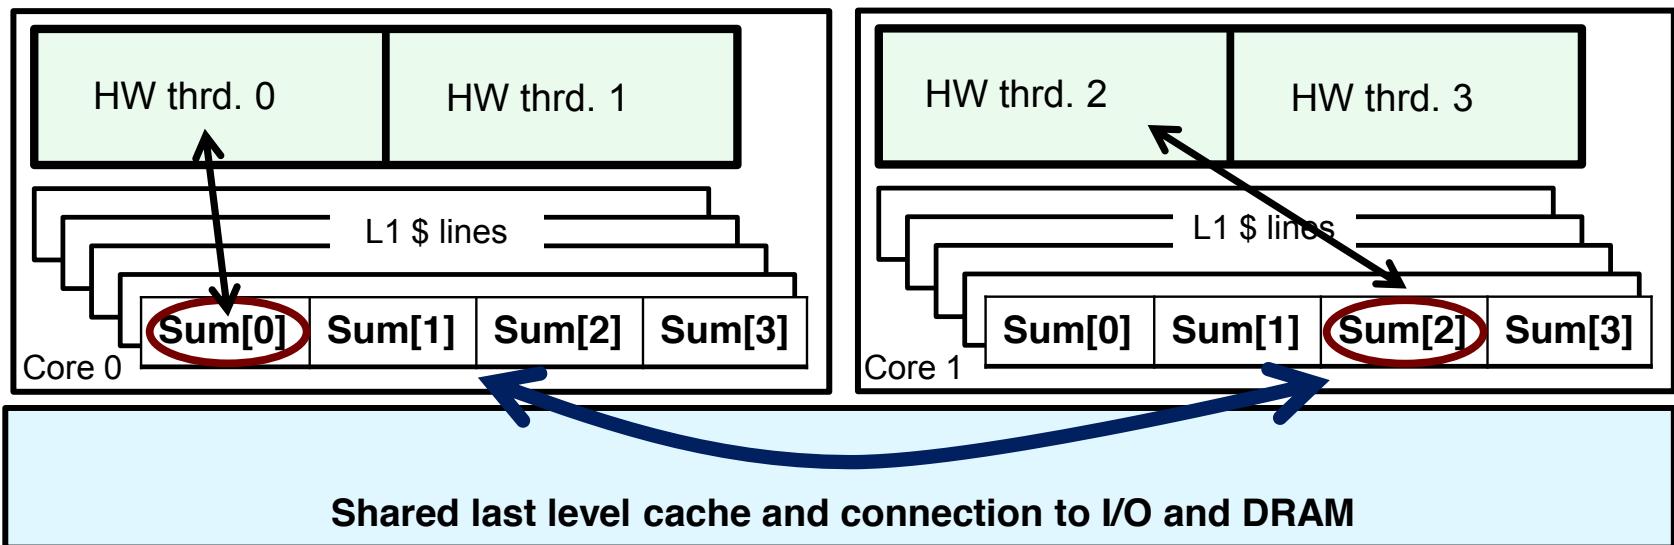
*Intel compiler (icpc) with no optimization on Apple OS X 10.7.3 with a dual core (four HW threads)
Intel® Core™ i5 processor at 1.7 Ghz and 4 Gbyte DDR3 memory at 1.333 Ghz



Why such poor scaling?

■ False sharing

- ⊕ If independent data elements happen to sit on the same cache line, each update will cause the cache lines to “slosh back and forth” between threads



- If you promote scalars to an array to support creation of an SPMD program, the array elements are contiguous in memory and hence share cache lines

- ⊕ Results in poor scalability

Solution: Pad arrays so elements you use are on distinct cache lines



Eliminating False Sharing by Padding sum

```
#include <omp.h>
static long num_steps = 100000; double step;
#define NUM_THREADS 2
#define PAD 8 // assume 64-byte L1 cache line size
void main () {
    int i, nthreads;
    double x, pi, sum[NUM_THREADS] [PAD];
    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int i, id, nthrds;
        double x;
        id = omp_get_thread_num();
        nthrds= omp_get_num_threads();
        if (id == 0) nthreads = nthrds;
        for (i=id, sum[id][0]=0.0; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            sum[id][0] += 4.0/(1.0+x*x);
        }
    }
    for(i=0, pi=0.0; i < nthreads; i++)
        pi += sum[i][0] * step;
}
```

Pad the array so each sum value is in a different cache line



Results*: Pi program (Padded Accumulator)

- Original Serial pi program with 100,000,000 steps ran in 1.83 seconds

```
#include <omp.h>
static long num_steps = 100000; double step;
#define NUM_THREADS 2
#define PAD 8 // assume 64-byte L1 cache line size
void main () {
    int i, nthreads;
    double x, pi, sum[NUM_THREADS] [PAD];
    step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
    #pragma omp parallel
    {
        int i, id, nthrds;
        double x;
        id = omp_get_thread_num();
        nthrds= omp_get_num_threads();
        if (id == 0) nthreads = nthrds;
        for (i=id, sum[id][0]=0.0; i< num_steps; i=i+nthrds) {
            x = (i+0.5)*step;
            sum[id][0] += 4.0/(1.0+x*x);
        }
    }
    for(i=0, pi=0.0; i < nthreads; i++)
        pi += sum[i][0] * step;
}
```

Threads	1 st SPMD	1 st SPMD (padded)		
1	1.86	0.98	1.86	0.98
2	1.03	0.89	1.01	0.9
3	1.08	0.56	0.69	0.88
4	0.97	0.47	0.53	0.86

*Intel compiler (icpc) with no optimization on Apple OS X 10.7.3 with a dual core (four HW threads)
Intel® Core™ i5 processor at 1.7 Ghz and 4 Gbyte DDR3 memory at 1.333 Ghz



Outline

- Introduction to OpenMP
- Compiler Directives
 - ❖ Creating Threads
 - ❖ Synchronization
 - ❖ Worksharing Constructs
 - ❖ More on Synchronization
 - ❖ Variable Scopes (shared/private)
- Run-time Libraries and Environment Variables



Synchronization

■ High level synchronization:

- critical
- atomic
- barrier
- ordered

Synchronization is used to impose order constraints and to protect access to shared data

■ Low level synchronization

- flush
- locks (both simple and nested)

Discussed later



Synchronization: critical

- Mutual exclusion: Only one thread at a time can enter a critical region

```
float res;  
#pragma omp parallel  
{  
    float B;    int i, id, nthrds;  
    id = omp_get_thread_num();  
    nthrds = omp_get_num_threads();  
    for(i=id;i<niters;i+nthrds) {  
        B = big_job(i);  
#pragma omp critical  
        res += consume(B);  
    }  
}
```

Threads wait
their turn -
only one at a
time calls
consume()



Synchronization: atomic (Basic Form)

- Atomic provides mutual exclusion but only applies to the update of a memory location (the update of `x` in the following example)
- The statement inside the atomic must be one of the following forms:

- `x binop= expr`
- `x++`
- `++x`
- `x--`
- `--x`

```
#pragma omp parallel
{
    double tmp, B;
    B = DOIT();
    tmp = big_ugly(B);
    #pragma omp atomic
    X += tmp;
}
```

- `x` is an lvalue of scalar type and `binop` is a non-overloaded built-in operator

Additional forms of atomic were added in OpenMP 3.1



critical Construct vs atomic Construct

critical

- ⊕ An OpenMP critical section is completely general
 - ◆ Incurring significant overhead every time a thread enters and exits the critical section
- ⊕ In OpenMP, all unnamed critical sections are considered identical
 - ◆ If one thread is in one [unnamed] critical section as above, no thread can enter any [unnamed] critical section
 - ◆ Naming a critical section: #pragma omp critical (name)

atomic

- ⊕ An atomic operation has much lower overhead. It relies on the hardware providing (say) an atomic increment operation
- ⊕ The upsides are that the overhead is much lower
- ⊕ The downsides are that you aren't guaranteed any particular set of atomic operations on any particular platform, and you could lose portability
 - ◆ However, the compiler should tell you if the particular atomic isn't supported



Some Caveats (1/2)

- You shouldn't mix the different types of mutual exclusion for a single critical section

- E.g.,

```
# pragma omp atomic      # pragma omp critical  
x += f(y);              x = g(x);
```

- The critical directive won't exclude the action executed by the atomic block

- There is no guarantee of fairness in mutual exclusion constructs

- A thread can be blocked forever

```
while(1) {  
    . . .  
#    pragma omp critical  
    x = g(my_rank);  
    . . .  
}
```



Some Caveats (2/2)

- It can be dangerous to “nest” mutual exclusion constructs

```
# pragma omp critical
y = f(x);
. . .
double f(double x) {
#   pragma omp critical
z = g(x); /* z is shared */
. . .
}
```

Guaranteed to deadlock

```
# pragma omp critical(one)
y = f(x);
. . .
double f(double x) {
#   pragma omp critical(two)
z = g(x); /* z is global */
. . .
}
```

- Naming critical sections might not help

Time	Thread <i>u</i>	Thread <i>v</i>
0	Enter crit. sect. one	Enter crit. sect. two
1	Attempt to enter two	Attempt to enter one
2	Block	Block



Using a critical section to remove impact of false sharing

```
#include <omp.h>
static long num_steps = 100000; double step;
#define NUM_THREADS 2
void main() {
    double x, pi; step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
    int i, id, nthrds; double x, sum;
    id = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    if (id == 0) nthrds = nthrds;
    for (i=id, sum=0.0; i < num_steps; i=i+nthrds) {
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
    }
#pragma omp critical
    pi += sum * step;
}
// end of parallel region
} // end of main
```

Create a scalar local to each thread to accumulate partial sums

No array, so no false sharing

sum goes “out of scope” beyond the parallel region ... so you must sum it in here. Must protect summation into pi in a critical region so updates don’t conflict



Results*: Pi program (CriticalSection)

- Original Serial pi program with 100,000,000 steps ran in 1.83 seconds

```
#include <omp.h>
static long num_steps = 100000; double step;
#define NUM_THREADS 2
void main() {
    double x, pi; step = 1.0/(double) num_steps;
    omp_set_num_threads(NUM_THREADS);
#pragma omp parallel
{
    int i, id, nthrds; double x, sum;
    id = omp_get_thread_num();
    nthrds = omp_get_num_threads();
    if (id == 0) nthrds = nthrds;
    for (i=id, sum=0.0; i < num_steps; i=i+nthrds) {
        x = (i+0.5)*step;
        sum += 4.0/(1.0+x*x);
    }
    #pragma omp critical
    pi += sum * step;
} // end of parallel region
} // end of main
```

Threads	1 st SPMD	1 st SPMD (padded)	SPMD (critical)
1	1.86	1.86	1.87
2	1.03	1.01	1.00
3	1.08	0.69	0.68
4	0.97	0.53	0.53

*Intel compiler (icpc) with no optimization on Apple OS X 10.7.3 with a dual core (four HW threads)
Intel® Core™ i5 processor at 1.7 Ghz and 4 Gbyte DDR3 memory at 1.333 Ghz



Outline

- Introduction to OpenMP
- Compiler Directives
 - ⊕ Creating Threads
 - ⊕ Synchronization
 - ⊕ Worksharing Constructs
 - ◆ Loop Constructs
 - ◆ Master/Single Constructs
 - ◆ Sections/Section Constructs
 - ⊕ More on Synchronization
 - ⊕ Variable Scopes (shared/private)
- Run-time Libraries and Environment Variables



SPMD vs Worksharing

- A parallel construct by itself creates an SPMD or “Single Program Multiple Data” program
 - ✿ i.e., each thread redundantly executes the same code
- How do you split up pathways through the code between threads within a team?
 - ✿ This is called worksharing
 - ◆ Loop construct
 - ◆ Master/single construct
 - ◆ Sections/section constructs
 - ◆ Task construct (available in OpenMP 3.0)

Discussed in the
next lecture



The Loop Worksharing Constructs

- The loop worksharing construct splits up loop iterations among the threads in a team

```
#pragma omp parallel
{
#pragma omp for
    for (I=0; I<N; I++) {
        NEAT_STUFF(I);
    }
}
```

- Loop construct name:
 - C/C++: for
 - Fortran: do

- The variable `I` is made “private” to each thread by default
- You could do this explicitly with a “`private(I)`” clause



Loop Worksharing Constructs: A Motivating Example

Sequential code

```
for(i=0; i<N; i++) {  
    a[i] = a[i] + b[i];  
}
```

OpenMP parallel region

```
#pragma omp parallel  
{  
    int id, Nthrds, istart, iend;  
    id = omp_get_thread_num();  
    Nthrds = omp_get_num_threads();  
    istart = id * N / Nthrds;  
    iend = (id+1) * N / Nthrds;  
    if (id == Nthrds-1)iend = N;  
    for(i=istart; i<iend; i++) {  
        a[i] = a[i] + b[i];  
    }  
}
```

OpenMP parallel region
and a worksharing for
construct

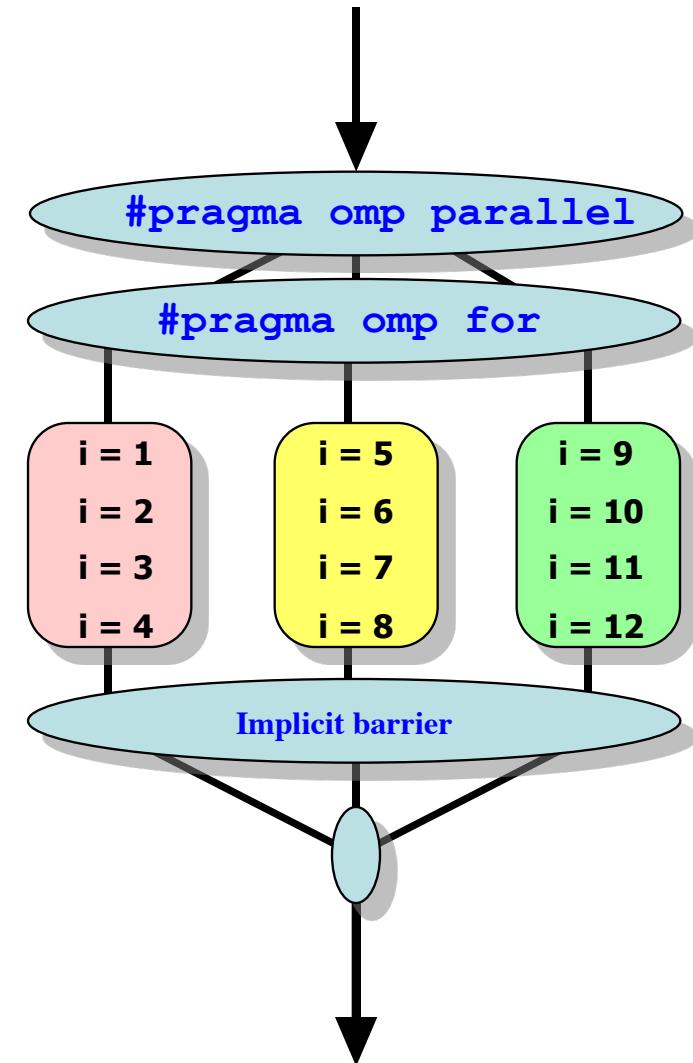
```
#pragma omp parallel  
#pragma omp for  
for(i=0; i<N; i++) {  
    a[i] = a[i] + b[i];  
}
```



omp for Construct

```
// assume N=12
#pragma omp parallel
#pragma omp for
for(i = 1; i < N+1; i++)
    c[i] = a[i] + b[i];
```

- Threads are assigned an independent set of iterations
- Threads must wait at the end of work-sharing construct



Loop Worksharing Constructs: schedule Clause (1/2)

- The schedule clause affects how loop iterations are mapped onto threads
 - ⊕ schedule(**static**[, chunk])
 - ◆ Deal-out blocks of iterations of size “chunk” to each thread
 - ⊕ schedule(**dynamic**[, chunk])
 - ◆ Each thread grabs “chunk” iterations off a queue until all iterations have been handled
 - ⊕ schedule(**guided**[, chunk])
 - ◆ Threads dynamically grab blocks of iterations. The size of the block starts large and shrinks down to size “chunk” as the calculation proceeds
 - ⊕ schedule(**runtime**)
 - ◆ Schedule and chunk size taken from the OMP_SCHEDULE environment variable (or the runtime library ... for OpenMP 3.0)
 - ⊕ schedule(**auto**)
 - ◆ Schedule is left up to the runtime to choose (does not have to be any of the above)

```
$ export OMP_SCHEDULE="static,1"
```



The static Schedule Type: Examples

■ Default schedule:

- `schedule(static, total_iterations/thread_count)`

■ `schedule(static, 1)`

Thread 0: 

Thread 1: 

Thread 2: 

■ `schedule(static, 2)`

Thread 0: 

Thread 1: 

Thread 2: 



The guided Schedule Type: Examples

Table 5.3 Assignment of Trapezoidal Rule Iterations 1–9999 using a guided Schedule with Two Threads

Thread	Chunk	Size of Chunk	Remaining Iterations
0	1–5000	5000	4999
1	5001–7500	2500	2499
1	7501–8750	1250	1249
1	8751–9375	625	624
0	9376–9687	312	312
1	9688–9843	156	156
0	9844–9921	78	78
1	9922–9960	39	39
1	9961–9980	20	19
1	9981–9990	10	9
1	9991–9995	5	4
0	9996–9997	2	2
1	9998–9998	1	1
0	9999–9999	1	0



Loop Worksharing Constructs: schedule Clause (1/2)

Schedule Clause	When To Use	
STATIC	Pre-determined and predictable by the programmer	Least work at runtime : scheduling done at compile-time
DYNAMIC	Unpredictable, highly variable work per iteration	Most work at runtime : complex scheduling logic used at run-time
GUIDED	Special case of dynamic to reduce scheduling overhead	
AUTO	When the runtime can “learn” from previous executions of the same loop	



Combined Parallel/Worksharing Construct

- OpenMP shortcut: Put the “parallel” and the worksharing directive on the same line

```
double res[MAX];  
int i;  
#pragma omp parallel  
{  
    #pragma omp for  
    for (i=0;i< MAX; i++) {  
        res[i] = huge();  
    }  
}
```

```
double res[MAX];  
int i;  
#pragma omp parallel for  
for (i=0;i< MAX; i++) {  
    res[i] = huge();  
}
```

These are equivalent



Working with Loops

Basic approach

- Find compute intensive loops
- Make the loop iterations independent
 - So they can safely execute in any order without loop-carried dependencies
- Place the appropriate OpenMP directive and test

```
int i, j, A[MAX];  
j = 5;  
for (i=0;i< MAX; i++) {  
    j += 2;  
    A[i] = big(j),  
}
```

Note: loop index “i”
is private by default

Remove loop-
carried dependence

```
int i, A[MAX];  
#pragma omp parallel for  
for (i=0;i< MAX; i++) {  
    int j = 5 + 2*(i+1);  
    A[i] = big(j);  
}
```



Nested Loops

- For perfectly nested rectangular loops we can parallelize multiple loops in the nest with the collapse clause:

```
#pragma omp parallel for collapse(2)
for (int i=0; i<N; i++) {
    for (int j=0; j<M; j++) {
        ...
    }
}
```

Number of loops to be parallelized, counting from the outside

- Will form a single loop of length $N \times M$ and then parallelize that
- Useful if N is $O(\text{no. of threads})$ so parallelizing the outer loop may complicate balancing the load



Reduction

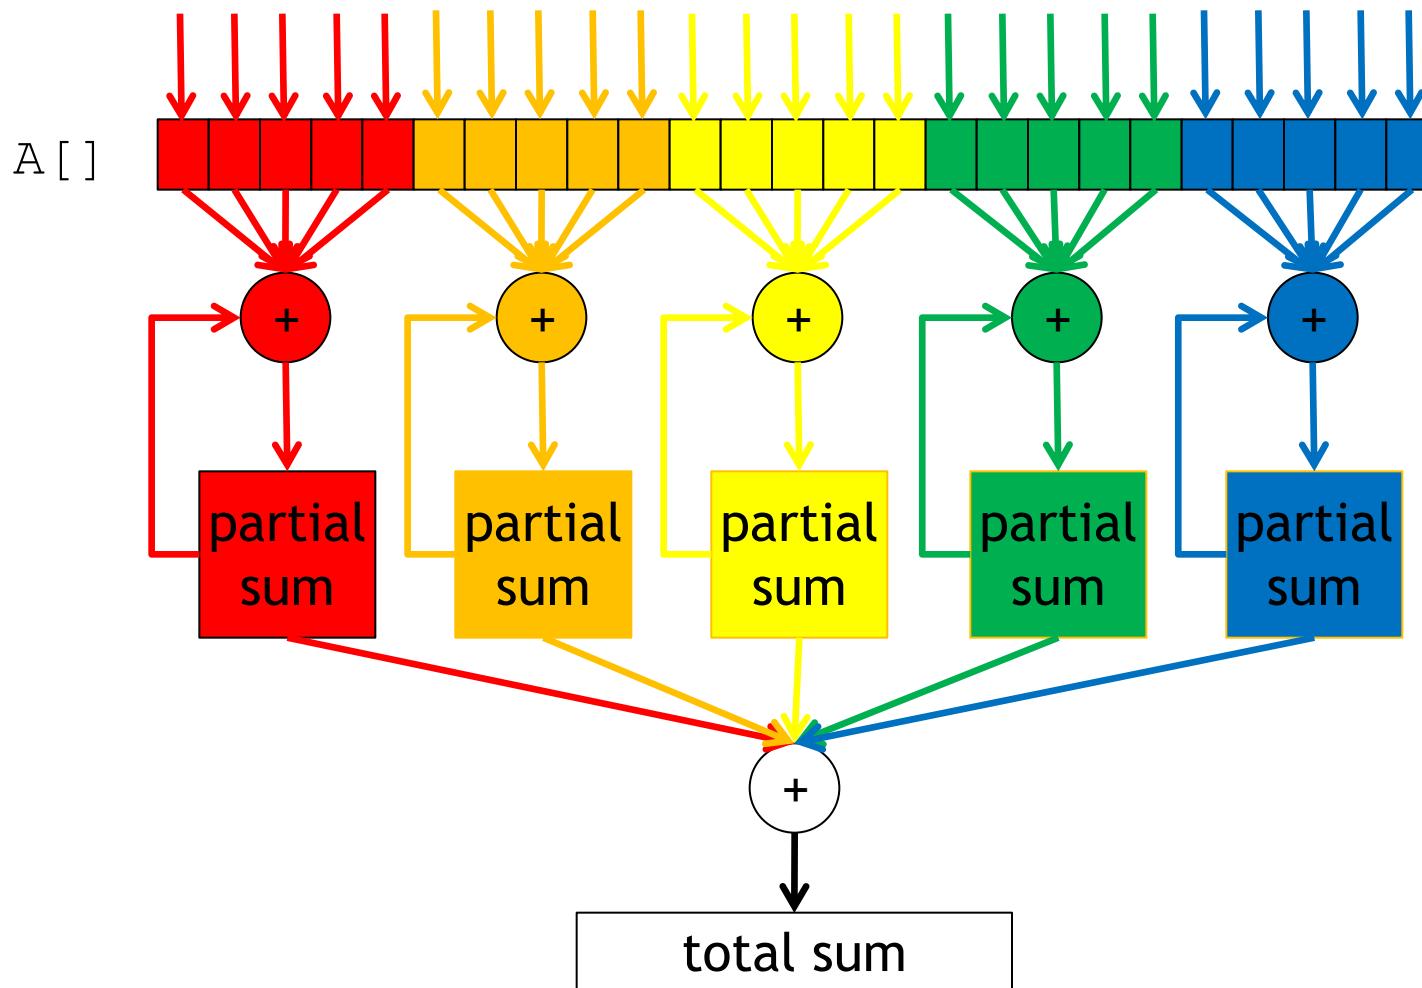
- How do we handle this case?

```
double ave=0.0, A[MAX];  
int i;  
for (i=0; i<MAX; i++) {  
    ave += A[i];  
}  
ave = ave/MAX;
```

- We are combining values into a single accumulation variable (ave)
 - ⊕ There is a true dependence between loop iterations that can't be trivially removed
- This is a very common situation (called a “reduction”)
- Support for reduction operations is included in most parallel programming environments



Reduction (Cont'd)



reduction Clause

- OpenMP reduction clause: `reduction (op : list)`
- Inside a parallel or a work-sharing construct:
 - ⊕ A local copy of each list variable is made and initialized depending on the “op” (**identity value**, e.g. 0 for “+”)
 - ⊕ Updates occur on the local copy
 - ⊕ Local copies are reduced into a single value and combined with the original global value
- The variables in “list” must be shared in the enclosing parallel region

```
double ave=0.0, A[MAX]; int i;  
#pragma omp parallel for reduction (+:ave)  
for (i=0; i<MAX; i++) {  
    ave += A[i];  
}  
ave = ave/MAX;
```



OpenMP: Reduction Operands/Initial-Values

- Many different associative operands can be used with reduction
- Initial values are the ones that make sense mathematically

Operator	Initial value
+	0
*	1
-	0

C/C++ only	
Operator	Initial value
&	~ 0
	0
^	0
&&	1
	0

Fortran Only	
Operator	Initial value
.AND.	.true.
.OR.	.false.
.NEQV.	.false.
.IEOR.	0
.IOR.	0
.IAND.	All bits on
.EQV.	.true.
MIN*	Largest pos. number
MAX*	Most neg. number



Example: Pi with a loop and a reduction

```
#include <omp.h>
static long num_steps = 100000;
double step;
void main() {
    int i; double pi, sum = 0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel {
        double x;
        #pragma omp for reduction(+:sum)
        for (i=0;i< num_steps; i++) {
            x = (i+0.5)*step;
            sum = sum + 4.0/(1.0+x*x);
        }
    }
    pi = step * sum;
}
```



Results*: Pi Program (omp for reduction)

- Original Serial pi program with 100,000,000 steps ran in 1.83 seconds

```
#include <omp.h>
static long num_steps = 100000;
double step;
void main() {
    int i; double x, pi, sum = 0.0;
    step = 1.0/(double) num_steps;
    #pragma omp parallel {
        double x;
        #pragma omp for reduction(+:sum)
        for (i=0;i< num_steps; i++) {
            x = (i+0.5)*step;
            sum = sum + 4.0/(1.0+x*x);
        }
    }
    pi = step * sum;
}
```

Threads	1 st SPMD	1 st SPMD (padded)	SPMD (critical)	Pi (Loop)
1	1.86	1.86	1.87	1.91
2	1.03	1.01	1.00	1.02
3	1.08	0.69	0.68	0.80
4	0.97	0.53	0.53	0.68

*Intel compiler (icpc) with no optimization on Apple OS X 10.7.3 with a dual core (four HW threads)
Intel® Core™ i5 processor at 1.7 Ghz and 4 Gbyte DDR3 memory at 1.333 Ghz



parallel for Caveats

- OpenMP compilers don't check for dependences among iterations in a parallel for loop
 - We need to worry about loop-carried dependences
- OpenMP will only parallelize for loops
 - It won't parallelize while or do-while loops
- OpenMP will only parallelize for loops for which the number of iterations can be determined
 - E.g., `for (; ;) { ... }` cannot be parallelized
 - OpenMP will only parallelize for loops that are in *canonical form*



for Loops in Canonical Form

- The variable `index` must have integer or pointer type (e.g., it can't be a float)
- The expressions `start`, `end`, and `incr` must have a compatible type
 - E.g., if `index` is a pointer, then `incr` must have integer type
- The expressions `start`, `end`, and `incr` must not change during execution of the loop
- During execution of the loop, the variable `index` can only be modified by the “increment expression” in the `for` statement

```
for ( index = start ; index >= end ; index += incr ) {  
    if ( index < end ) {  
        index++  
        ++index  
    } else if ( index <= end ) {  
        index--  
        --index  
    } else if ( index > end ) {  
        index += incr  
        index -= incr  
        index = index + incr  
        index = incr + index  
        index = index - incr  
    }  
}
```



Outline

- Introduction to OpenMP
- Compiler Directives
 - ⊕ Creating Threads
 - ⊕ Synchronization
 - ⊕ Worksharing Constructs
 - ◆ Loop Constructs
 - ◆ **Master/Single Constructs**
 - ◆ Sections/Section Constructs
 - ⊕ More on Synchronization
 - ⊕ Variable Scopes (shared/private)
- Run-time Libraries and Environment Variables



master Construct

- The master construct denotes a structured block that is only executed by the master thread
- The other threads just skip it (no synchronization is implied)

```
#pragma omp parallel
{
    do_many_things();
#pragma omp master
{
    exchange_boundaries();
}
#pragma omp barrier
    do_many_other_things();
}
```



single Worksharing Construct

- The single construct denotes a block of code that is executed by only one thread (not necessarily the master thread)
- A barrier is implied at the end of the single block (can remove the barrier with a nowait clause)

```
#pragma omp parallel
{
    do_many_things();
#pragma omp single
{
    exchange_boundaries();
}
    do_many_other_things();
}
```



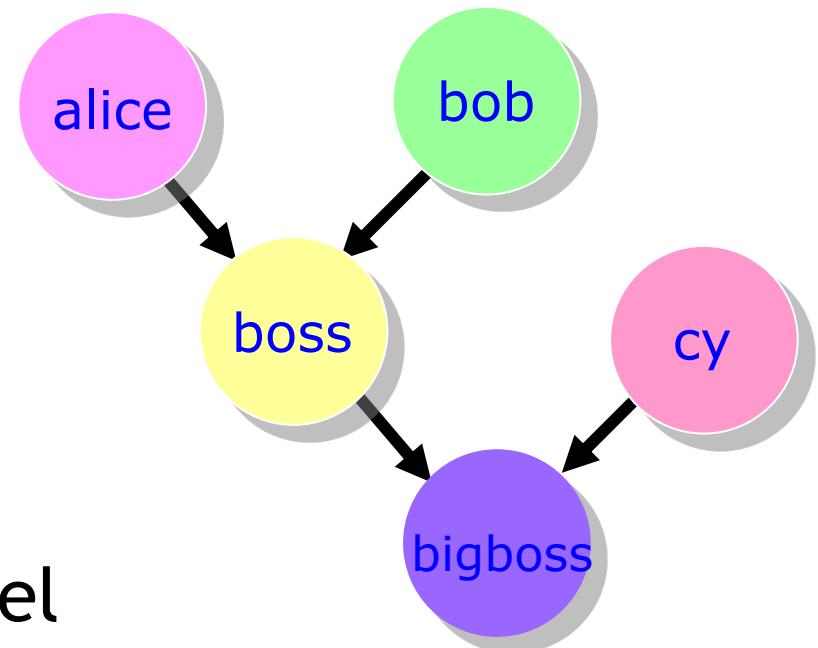
Outline

- Introduction to OpenMP
- Compiler Directives
 - ⊕ Creating Threads
 - ⊕ Synchronization
 - ⊕ Worksharing Constructs
 - ◆ Loop Constructs
 - ◆ Master/Single Constructs
 - ◆ **Sections/Section Constructs**
 - ⊕ More on Synchronization
 - ⊕ Variable Scopes (shared/private)
- Run-time Libraries and Environment Variables



Function-Level Parallelism

```
a = alice();  
b = bob();  
s = boss(a, b);  
c = cy();  
printf ("%6.2f\n", bigboss(s, c));
```



alice, bob, and cy
can be computed in parallel



omp sections/section

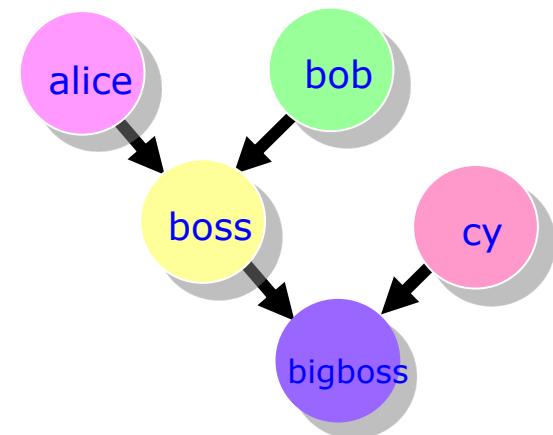
- #pragma omp sections
 - ❖ Must be inside a parallel region
 - ❖ Precedes a code block containing N sub-blocks of code that may be executed concurrently by N threads
 - ❖ Encompasses each omp section
- #pragma omp section
 - ❖ Precedes each sub-block of code within the encompassing block described above
 - ❖ Enclosed program segments are distributed for parallel execution among available threads



Function-Level Parallelism Using `omp sections`

```
#pragma omp parallel sections
{
    #pragma omp section
        double a = alice();
    #pragma omp section
        double b = bob();
    #pragma omp section
        double c = cy();
}

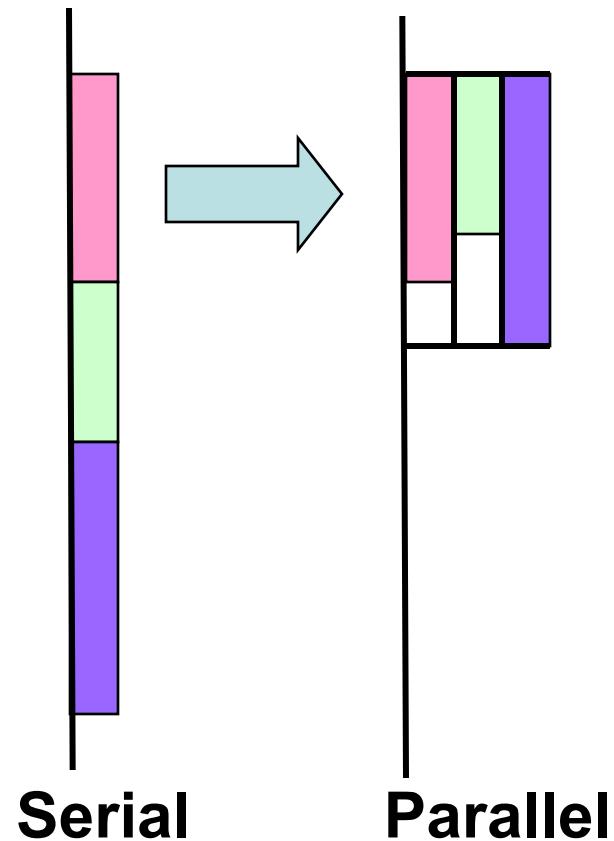
double s = boss(a, b);
printf ("%6.2f\n", bigboss(s, c));
```



Advantages of Parallel Sections

- Independent sections of code can execute concurrently - reduce execution time

```
#pragma omp parallel sections
{
#pragma omp section
    phase1();
#pragma omp section
    phase2();
#pragma omp section
    phase3();
}
```



sections Worksharing Construct

- The Sections worksharing construct gives a different structured block to each thread

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
            x_calculation();
        #pragma omp section
            y_calculation();
        #pragma omp section
            z_calculation();
    }
}
```

By default, there is a barrier at the end of the “omp sections”.
Use the “nowait” clause to turn off the barrier.



Outline

- Introduction to OpenMP
- Compiler Directives
 - ❖ Creating Threads
 - ❖ Synchronization
 - ❖ Worksharing Constructs
 - ❖ More on Synchronization
 - ❖ Variable Scopes (shared/private)
- Run-time Libraries and Environment Variables



Synchronization: barrier

- Barrier: Each thread waits until all threads arrive

```
#pragma omp parallel shared (A, B, C) private(id) {  
    id = omp_get_thread_num();  
    A[id] = big_calc1(id);  
#pragma omp barrier  
#pragma omp for  
    for(i=0;i<N;i++) { C[i] = big_calc3(i,A); }  
#pragma omp for nowait  
    for(i=0;i<N;i++) { B[i]=big_calc2(C, i); }  
    A[id] = big_calc4(id);  
}  
implicit barrier at the end of a parallel region  
implicit barrier at the end of a for worksharing construct  
no implicit barrier due to nowait
```



Synchronization: ordered

- The ordered region executes in the sequential order

```
#pragma omp parallel private (tmp)
#pragma omp for ordered reduction(+:res)
for (I=0; I<N; I++) {
    tmp = NEAT_STUFF(I);
#pragma omp ordered
    res += consum(tmp);
}
```



Outline

- Introduction to OpenMP
- Compiler Directives
 - ❖ Creating Threads
 - ❖ Synchronization
 - ❖ Worksharing Constructs
 - ❖ More on Synchronization
 - ❖ **Variable Scopes (shared/private)**
- Run-time Libraries and Environment Variables



Default Storage Attributes

- Shared memory programming model:
 - ✿ Most variables are shared by default
- Global variables are SHARED among threads
 - ✿ Fortran: COMMON blocks, SAVE variables, MODULE variables
 - ✿ C: File scope variables, static variables
 - ✿ Both: dynamically allocated memory (malloc, new)
- But not everything is shared...
 - ✿ Stack variables in subprograms(Fortran) or functions(C) called from parallel regions are PRIVATE
 - ✿ Automatic variables within a statement block are PRIVATE



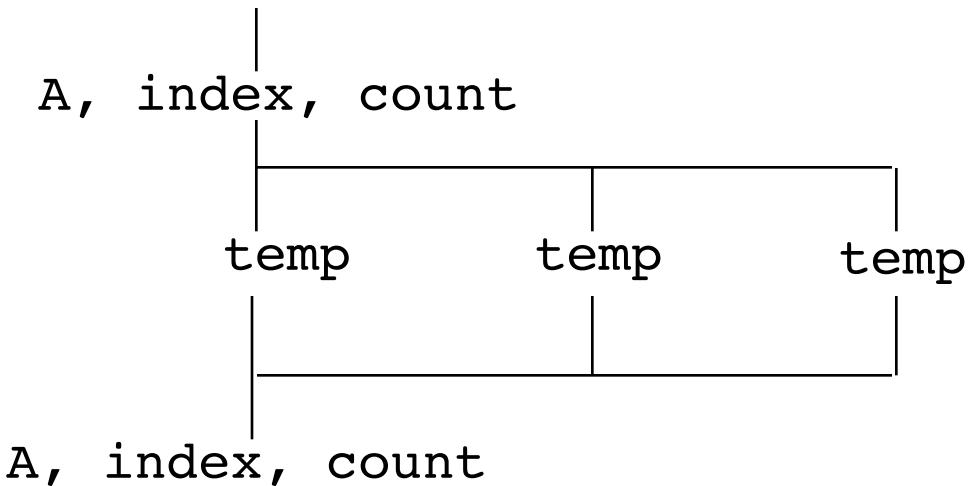
Data Sharing: Examples

```
double A[10];
int main() {
    int index[10];
#pragma omp parallel
    work(index);
    printf("%d\n", index[0]);
}
```

A, index and count
are shared by all threads.

temp is local to each
thread

```
extern double A[10];
void work(int *index) {
    double temp[10];
    static int count;
    ...
}
```



Data Sharing: Changing Storage Attributes

- One can selectively change storage attributes for constructs using the following clauses*
 - shared
 - private
 - firstprivate
- The final value of a private inside a parallel loop can be transmitted to the shared variable outside the loop with:
 - lastprivate
- The default attributes can be overridden with:
 - default (private | shared | none)

All the clauses on this page apply to the OpenMP construct NOT to the entire region

*All data clauses apply to parallel constructs and worksharing constructs except “shared” which only applies to parallel constructs



Data Sharing: private Clause

- `private(var)` creates a new local copy of var for each thread
 - The value is uninitialized
 - In OpenMP 2.5 the value of the private variable is undefined after the region

```
void wrong() {  
    int tmp = 0;  
#pragma omp parallel for private(tmp)  
    for (int j = 0; j < 1000; ++j)  
        tmp += j;  
    printf("%d\n", tmp);  
}
```

tmp was not initialized

tmp: unspecified in 2.5, 0 in 3.0 (implementations may reference the original variable or a copy. A dangerous programming practice!)



Data Sharing: firstprivate Clause

- `firstprivate` is a special case of `private`
 - Variables initialized from shared variable
 - C++ objects are copy-constructed

```
void useless() {  
    int tmp = 0;  
#pragma omp parallel for firstprivate(tmp)  
    for (int j = 0; j < 1000; ++j)  
        tmp += j; ←  
    printf("%d\n", tmp);  
}
```

Each thread gets its own `tmp` with an initial value of 0

tmp: 0 in 3.0, unspecified in 2.5



Data Sharing: lastprivate Clause

- lastprivate is a special case of private
 - ✿ Variables update shared variable using value from last iteration
 - ✿ C++ objects are updated as if by assignment

```
void closer() {  
    int tmp = 0;  
#pragma omp parallel for firstprivate(tmp) \  
    lastprivate(tmp)  
    for (int j = 0; j < 1000; ++j)  
        tmp += j;  
    printf("%d\n", tmp);  
}
```

Each thread gets its own tmp with
an initial value of 0

tmp is defined as its value at the “last
sequential” iteration (i.e., for j=999)



Data Sharing: default Clause

- Note that the default storage attribute is default (shared) (so no need to use it)
 - ✿ Exception: #pragma omp task
- To change default: default (private)
 - ✿ Each variable in the construct is made private as if specified in a private clause
 - ✿ Mostly saves typing
- default (none) : no default for variables in static extent
 - ✿ Must list storage attribute for each variable in static extent
 - ✿ Good programming practice!

Only the Fortran API supports default (private)
C/C++ only has default (shared) or default (none)



Example: Pi Program with Minimal Changes

```
#include <omp.h>
static long num_steps = 100000;
double step;
void main() {
    int i; double x, pi, sum = 0.0;
    step = 1.0/(double) num_steps;
#pragma omp parallel for private(x) \
reduction(+:sum)
    for (i=0; i<num_steps; i++) {
        x = (i+0.5)*step;
        sum = sum + 4.0/(1.0+x*x);
    }
    pi = step * sum;
}
```

**i is private
by default**

For good OpenMP implementations,
reduction is more scalable than critical

Note: we created a parallel program
without changing any executable code
and by adding 2 simple lines of text!



Outline

- Introduction to OpenMP
- Compiler Directives
 - ❖ Creating Threads
 - ❖ Synchronization
 - ❖ Worksharing Constructs
 - ❖ More on Synchronization
 - ❖ Variable Scopes (shared/private)
- Run-time Libraries and Environment Variables



Synchronization: Lock Routines

Simple Lock routines:

- A simple lock is available if it is unset
 - ◆ `omp_init_lock()`, `omp_set_lock()`,
`omp_unset_lock()`, `omp_test_lock()`,
`omp_destroy_lock()`

A lock implies a memory fence (a “flush”) of all thread visible variables

Nested Locks

- A nested lock is available if it is unset or if it is set but owned by the thread executing the nested lock function
 - ◆ `omp_init_nest_lock()`, `omp_set_nest_lock()`,
`omp_unset_nest_lock()`, `omp_test_nest_lock()`,
`omp_destroy_nest_lock()`

Note: a thread always accesses the most recent copy of the lock, so you don't need to use a flush on the lock variable.



Example: Simple Locks

■ Protect resources with locks

```
omp_lock_t lck;  
omp_init_lock(&lck);  
#pragma omp parallel private (tmp, id)  
{  
    id = omp_get_thread_num();  
    tmp = do_lots_of_work(id);  
    omp_set_lock(&lck);           ← Wait here for your turn  
    printf("%d %d", id, tmp);  
    omp_unset_lock(&lck);          ← Release the lock so the  
                                    next thread gets a turn  
}  
omp_destroy_lock(&lck);          ← Free-up storage when done
```



Runtime Library Routines

■ Runtime environment routines:

- Modify/check the number of threads
 - ◆ `omp_set_num_threads()`, `omp_get_num_threads()`,
`omp_get_thread_num()`, `omp_get_max_threads()`
 - Are we in an active parallel region?
 - ◆ `omp_in_parallel()`
 - Do you want the system to dynamically vary the number of threads from one parallel construct to another?
 - ◆ `omp_set_dynamic()`, `omp_get_dynamic()`
 - How many processors in the system?
 - ◆ `omp_num_procs()`
- ## ■ ... plus a few less commonly used routines



Runtime Library Routines: Examples

- To use a known, fixed number of threads in a program,
(1) tell the system that you don't want dynamic adjustment of the
number of threads, (2) set the number of threads, then (3) save the
number you got

```
#include <omp.h>
void main() {
    int num_threads;
    omp_set_dynamic(0);
    omp_set_num_threads(omp_get_num_procs());
#pragma omp parallel
{
    int id = omp_get_thread_num();
#pragma omp single
    num_threads = omp_get_num_threads();
    do_lots_of_stuff(id);
}
}
```

Disable dynamic adjustment of the number of threads

Request as many threads as you have processors

Protect this op since memory stores are not atomic

Even in this case, the system may give you fewer threads than requested. If the precise # of threads matters, test for it and respond accordingly.



Environment Variables

- Set the default number of threads to use
 - ⊕ OMP_NUM_THREADS *int_literal*
- OpenMP added an environment variable to control the size of child threads' stack
 - ⊕ OMP_STACKSIZE
- Also added an environment variable to hint to runtime how to treat idle threads
 - ⊕ OMP_WAIT_POLICY
 - ⊕ ACTIVE keep threads alive at barriers/locks
 - ⊕ PASSIVE try to release processor at barriers/locks
- Control how “omp for schedule(runtime)” loop iterations are scheduled
 - ⊕ OMP_SCHEDULE “*schedule[, chunk_size]*”
- ... Plus several less commonly used environment variables



Summary

- OpenMP is one of the simplest APIs available for programming shared memory machines
 - This simplicity means you can focus on mastering the key design patterns and applying them to your own problems
- We covered the following essential parallel programming design patterns:
 - Fork join
 - SPMD
 - Loop level parallelism
- Next steps?
 - Let's consider some of the newer and more advanced features of OpenMP



References

- Tim Mattson, Michael Wrinn, and Mark Bull, “A “Hands-on” Introduction to OpenMP”
- Tim Mattson, “Shared Memory Programming with OpenMP - Basics”, 2012 Short Course on Parallel Programming, Berkeley
- Tim Mattson, “More about OpenMP- New Features” , 2012 Short Course on Parallel Programming, Berkeley
- Intel Software College, “Programming with OpenMP”
- Gabriel Mateescu, “Writing and tuning OpenMP programs on distributed shared memory platforms”
- UCB CS267 Course, “Applications of Parallel Computers”

