# Parallel Programming

## Hadoop and MapReduce Programming

Professor Yi-Ping You (游逸平)

Department of Computer Science

http://www.cs.nctu.edu.tw/~ypyou/

# Acknowledgement

- This set of slides is adapted from lectures by
    - Prof. Jimmy Lin (University of Maryland), "Introduction to Cloud Computing"
    - Prof. Anand Rajaraman (Stanford University), "Data Mining"
    - Prof. Adriana Iamnitchi (University of South Florida), "Basics of Parallel and Distributed Systems"

# Outline

- Introduction to MapReduce
  - Google File System
- The Hadoop Framework

# What is MapReduce?

- Data-parallel programming model for clusters of commodity machines

- Pioneered by Google (published in OSDI'04)
  - Processes 20 PB of data per day
    - Index building for Google Search
    - Article clustering for Google News
    - Statistical machine translation

- Popularized by Apache Hadoop project
  - Used by Yahoo!, Facebook, Amazon, ...
    - Index building for Yahoo! Search
    - Spam detection for Yahoo! Mail
    - Ad optimization for Facebook

# Commodity Clusters

- ☐ Web data sets can be very large
  - ■ Tens to hundreds of terabytes
- ☐ Cannot mine on a single server (why?)
- ☐ Standard architecture emerging:
  - ■ Cluster of commodity Linux nodes
  - ■ Gigabit ethernet interconnect
- ☐ How to organize computations on this architecture?
  - ■ Mask issues such as hardware failure

# Cluster Architecture

2-10 Gbps backbone between racks

1 Gbps between any pair of nodes in a rack

Switch

Switch

Switch

| CPU | CPU | CPU | CPU |
|-----|-----|-----|-----|
| Mem | Mem | Mem | Mem |
| Disk | Disk | Disk | Disk |

...

...

Each rack contains 16-64 nodes

# Stable storage

- ☐ First order problem: if nodes can fail, how can we store data persistently?
- ☐ Answer: Distributed File System
  - ■ Provides global file namespace
  - ■ Google GFS; Hadoop HDFS; Kosmix KFS
- ☐ Typical usage pattern
  - ■ Huge files (100s of GB to TB)
  - ■ Data is rarely updated in place
  - ■ Reads and appends are common

# Distributed File System

☐ Chunk Servers
- ■ a.k.a. Data Nodes in HDFS
- ■ File is split into contiguous chunks
- ■ Typically each chunk is 16-64MB
- ■ Each chunk replicated (usually 2x or 3x)
- ■ Try to keep replicas in different racks

☐ Master node
- ■ a.k.a. Name Node in HDFS
- ■ Stores metadata
- ■ Might be replicated

☐ Client library for file access
- ■ Talks to master to find chunk servers
- ■ Connects directly to chunkservers to access data

# Motivation for MapReduce

- ☐ Large-Scale Data Processing
  - ◼ Want to use 1000s of CPUs
    - ☐ But don't want hassle of **managing** things

- ☐ MapReduce Architecture provides
  - ◼ Automatic parallelization & distribution
  - ◼ Fault tolerance
  - ◼ I/O scheduling
  - ◼ Monitoring & status updates

# MapReduce Goals

- Cloud Environment:
  - Commodity nodes (cheap, but unreliable)
  - Commodity network (low bandwidth)
  - Automatic fault-tolerance (fewer admins)
- Scalability to large data volumes:
  - Scan 100 TB on 1 node @ 50 MB/s = 24 days
  - Scan on 1000-node cluster = 35 minutes

# What is Map/Reduce

- ☐ Map/Reduce
    - ■ Programming model from LISP
    - ■ (and other functional languages)

- ☐ Many problems can be phrased this way

- ☐ Easy to distribute across nodes
- ☐ Nice retry/failure semantics

# Map in LISP (Scheme)

- (map ***f list [list$_2$ list$_3$ ...]***)

```
foreach (x in list) {
  f(x);
}
```

*Unary operator*

- (map square '(1 2 3 4))
  - (1 4 9 16)

# Map

☐ Map is a higher-order function

☐ How map works:

■ Function is applied to every element in a list

■ Result is a new list

☐ No limit to map parallelization since maps are independent

# Reduce in LISP (Scheme)

☐ (reduce **f id list**)

*Binary operator*

☐ (reduce + 0  '(1 4 9 16))

  ■ (+ 16 (+ 9 (+ 4 (+ 1 0)) ) )

  ■ 30

☐ (reduce + 0
        (map square  '($l_1$ $l_2$)))

# Reduce

☐ Reduce is also a higher-order function

☐ How reduce works:

- ◼ Accumulator set to initial value
- ◼ Function applied to list element and the accumulator
- ◼ Result stored in the accumulator
- ◼ Repeated for every item in the list
- ◼ Result is the final value in the accumulator

# Reduce



☐ We can reorder folding if the reduce function is commutative and associative

# Typical Problem

- ☐ Iterate over a large number of records

*Map*

- ☐ Extract something of interest from each
- ☐ Shuffle and sort intermediate results
- ☐ Aggregate intermediate results

*Reduce*

- ☐ Generate final output

**Key idea:** provide an abstraction at the point of these two operations

# MapReduce

☐ Programmers specify two functions:

Map function:

$$(K_{in}, V_{in}) \ ➡ \ list<(K_{inter}, V_{inter})>$$

Reduce function:

$$(K_{inter}, list<V_{inter}>) ➡ list<(K_{out}, V_{out})>$$

☐ Usually, programmers also specify:

**partition** (k', number of partitions ) → partition for k'

■ Often a simple hash of the key, e.g. hash(k') mod n

■ Allows reduce operations for different keys in parallel

# It's just divide and conquer
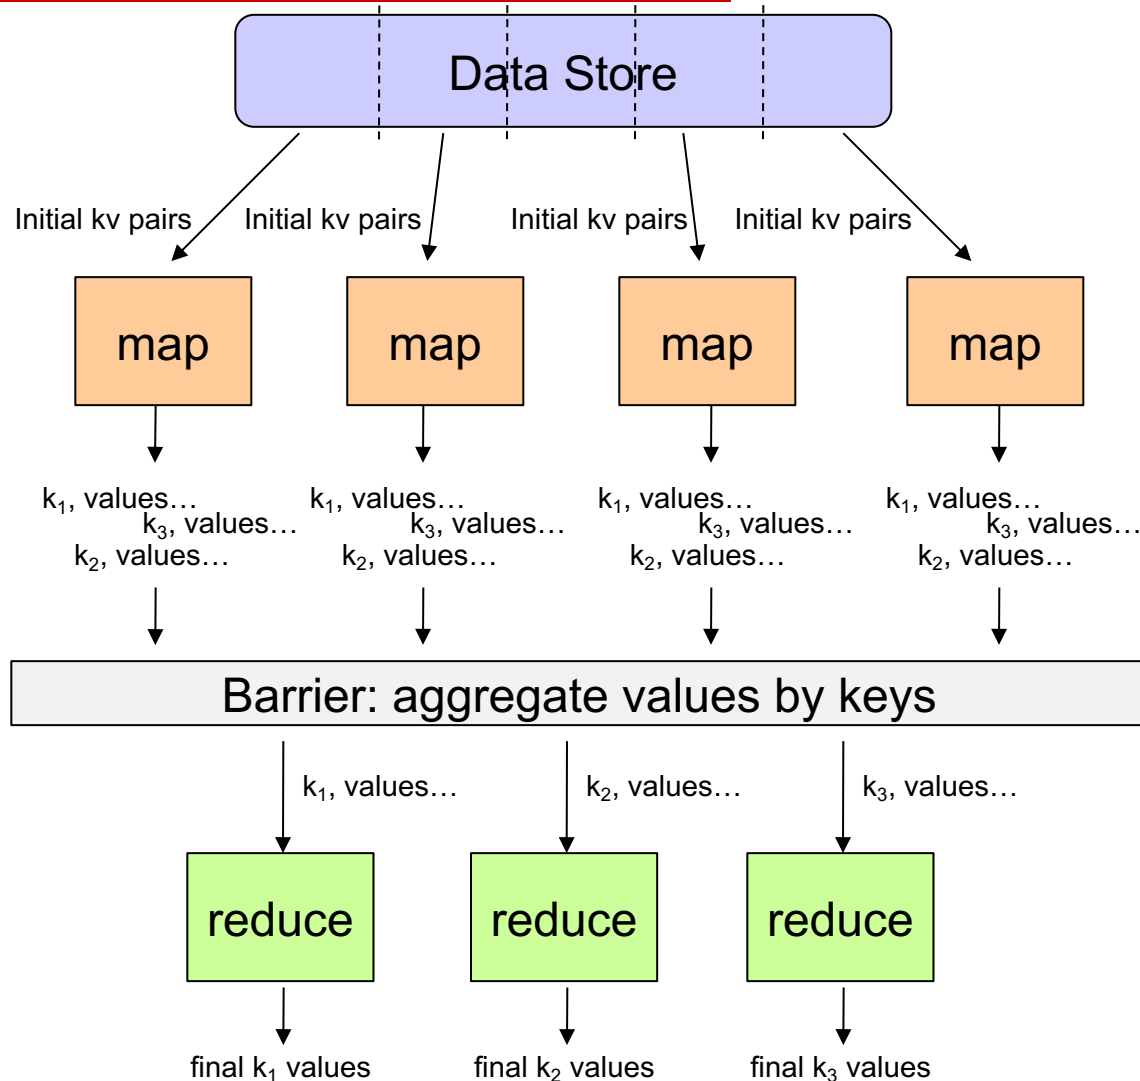
# "Hello World": Word Count

- ☐ We have a large file of words

- ☐ Count the number of times each distinct word appears in the file

- ☐ *Sample application*: analyze web server logs to find popular URLs

# Word Count (2)

- ☐ Case 1: Entire file fits in memory
- ☐ Case 2: File too large for mem, but all <word, count> pairs fit in mem
- ☐ Case 3: File on disk, too many distinct words to fit in memory
  - ■ `xargs -n1 < datafile | sort | uniq -c`

# Word Count (3)

- ☐ To make it slightly harder, suppose we have a large corpus of documents
- ☐ Count the number of times each distinct word occurs in the corpus

  **`words(docs/*) | sort | uniq -c`**

  where **`words`** takes a file and outputs the words in it, one to a line

- ☐ The above captures the essence of MapReduce
  - ▪ Great thing is that it is naturally parallelizable

# MapReduce

☐ Input: a set of key/value pairs
☐ User supplies two functions:
  ■ map(k,v) → list(k1,v1)
  ■ reduce(k1, list(v1)) → v2
☐ (k1,v1) is an intermediate key/value pair
☐ Output is the set of (k1,v2) pairs

# Word Count using MapReduce

```
map(key, value):
// key: document name; value: text of document
    for each word w in value:
        emit(w, 1)



reduce(key, values):
// key: a word; values: an iterator over counts
        result = 0
        for each count v in values:
                result += v
        emit(key,result)
```

# Count, Illustrated

map(key=url, val=contents):

For each word *w* in contents, emit (w,"1")

reduce(key=word, values=uniq_counts):

Sum all "1" s in values list
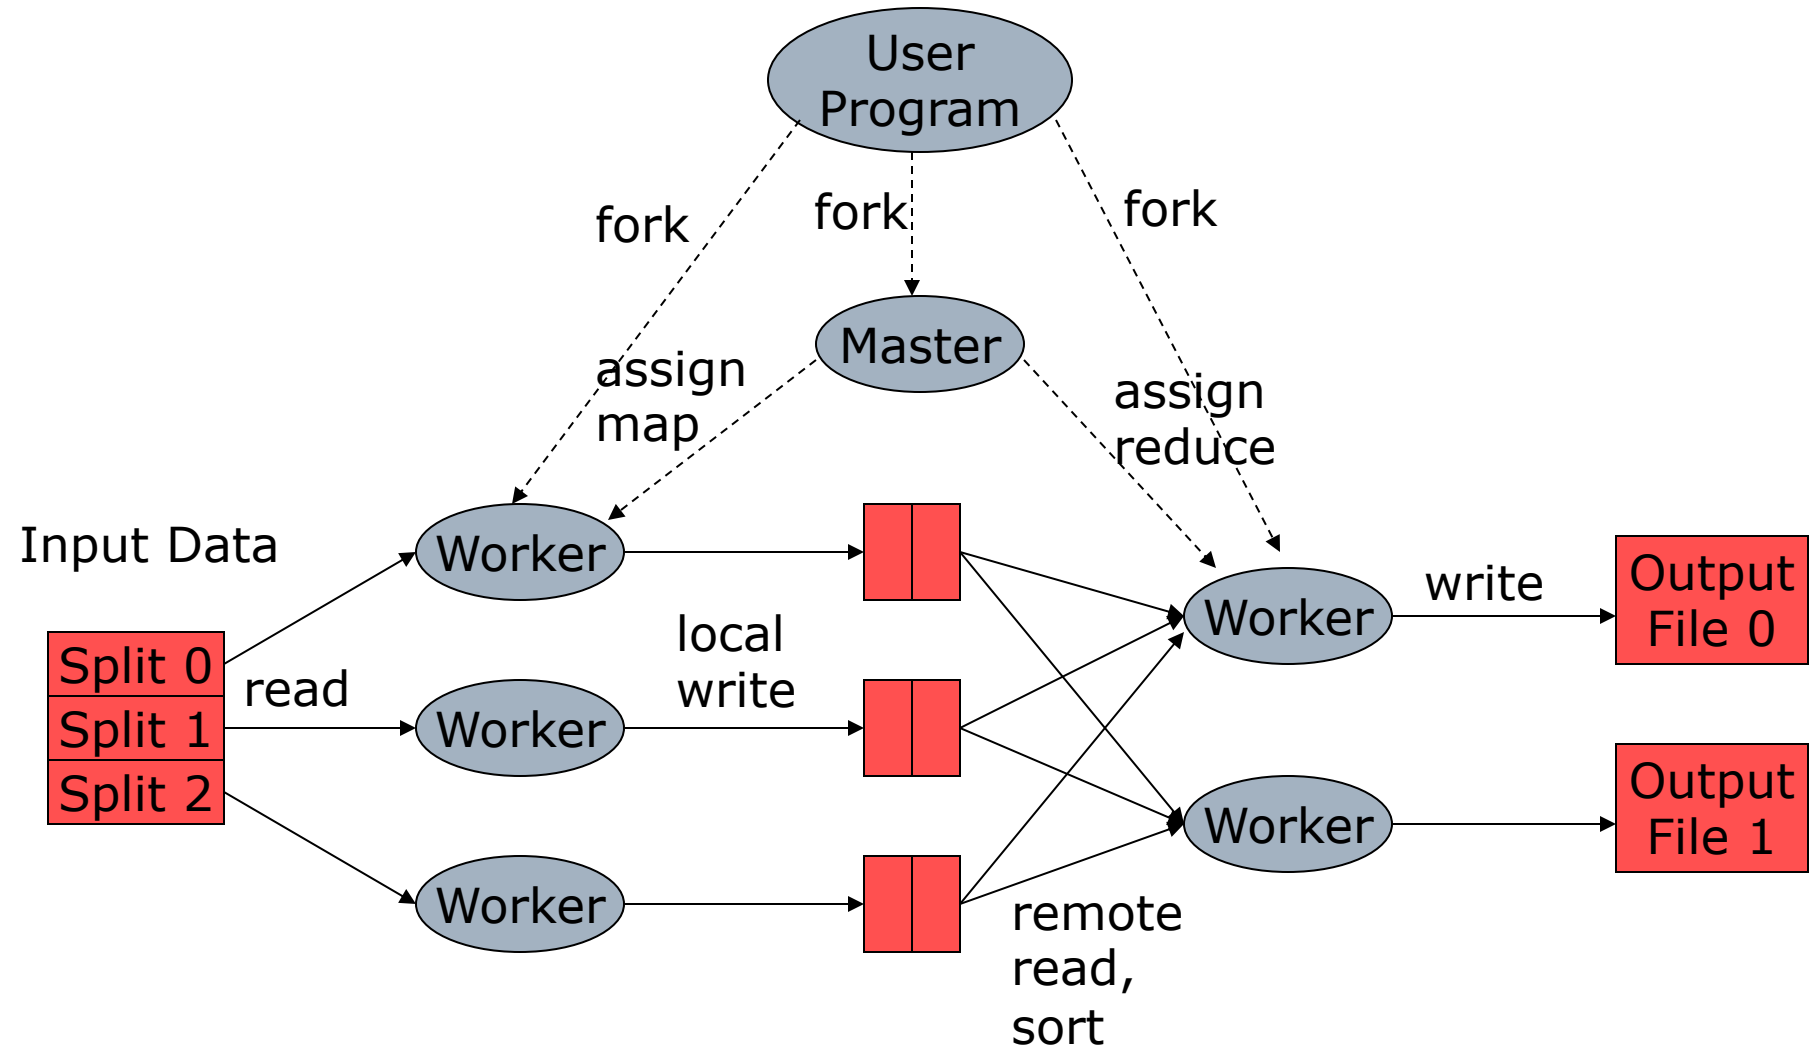
Emit result "(word, sum)"

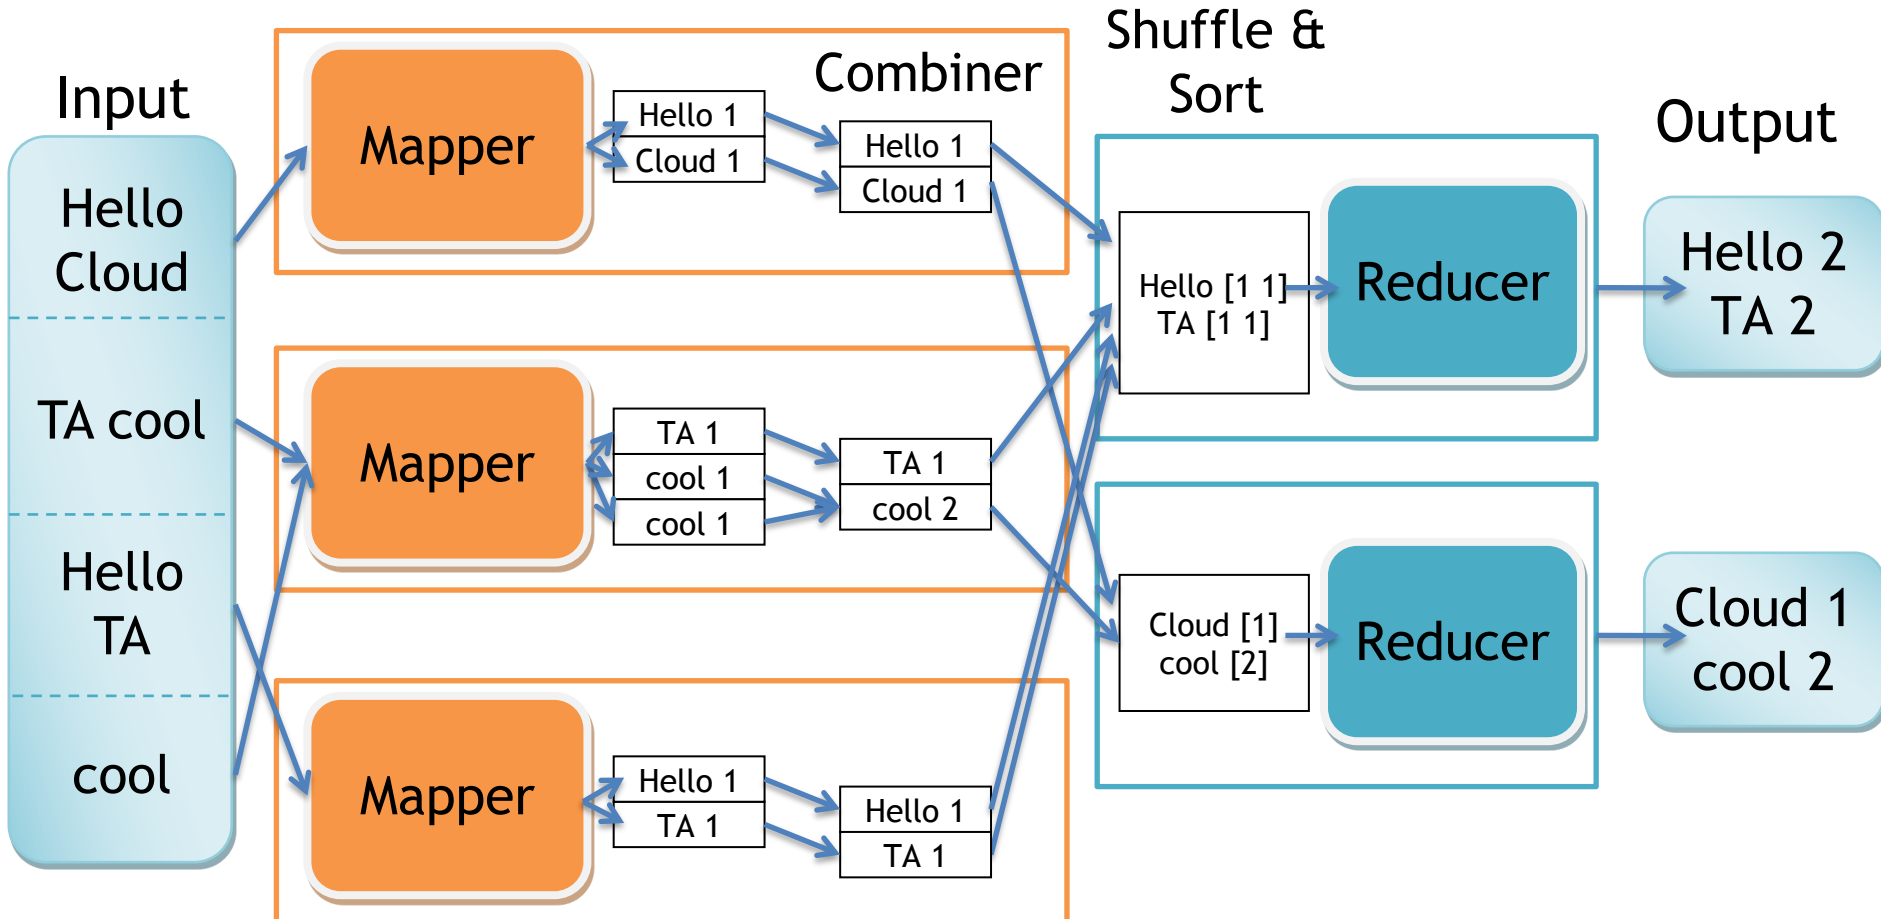| see bob run | see | 1 |   | bob | 1 |
|:---|:---|:---|:---|:---|:---|
| see spot throw | bob | 1 |   | run | 1 |
|   | run | 1 |   | see | 2 |
|   | see | 1 |   | spot | 1 |
|   | spot | 1 |   | throw | 1 |
|   | throw | 1 |   |   |   |

# Behind the scenes…

# Bandwidth Optimization

☐ Issue: large number of key-value pairs

☐ Solution: use "Combiner" functions

- ■ Executed on same machine as mapper
- ■ Results in a "mini-reduce" right after the map phase
- ■ Reduces key-value pairs to save bandwidth

# Example: Word Count

```
def map(line_num, line):
    foreach word in line.split():
        output(word, 1)
```

```
def reduce(word, counts):
    output(word, sum(counts))
```



Input

Hello Cloud

TA cool

Hello TA

cool

Mapper

Mapper

Mapper

Combiner

Hello 1
Cloud 1

Hello 1
Cloud 1

TA 1
cool 1
cool 1

TA 1
cool 2

Hello 1
TA 1

Hello 1
TA 1

Shuffle & Sort

Hello [1 1]
TA [1 1]

Cloud [1]
cool [2]

Reducer

Reducer

Output

Hello 2
TA 2

Cloud 1
cool 2

# Combiners

- ☐ Often a map task will produce many pairs of the form (k,v1), (k,v2), … for the same key k
  - ■ E.g., popular words in Word Count
- ☐ Can save network time by pre-aggregating at mapper
  - ■ combine(k1, list(v1)) → v2
  - ■ Usually same as reduce function
- ☐ Works only if reduce function is commutative and associative

# Skew Problem

☐ **Issue:** reduce is only as fast as the slowest map

☐ **Solution:** redundantly execute map operations, use results of first to finish

  ◼ Addresses hardware problems...

  ◼ But not issues related to inherent distribution of data

# Model is Widely Applicable
## MapReduce Programs In Google Source Tree



Example uses:

| | | |
|---|---|---|
| distributed grep | distributed sort | web link-graph reversal |
| term-vector / host | web access log stats | inverted index construction |
| document clustering | machine learning | statistical machine translation |
| ... | ... | ... |

# Example: Reverse Web-Link Graph

☐ Find all pages that link to a certain page
☐ Map Function
  ◼ Outputs *<target, source>* pairs for each link to a *target* URL found in a *source* page
  ◼ For each page we know what pages it links to
☐ Reduce Function
  ◼ Concatenates the list of all *source* URLs associated with a given *target* URL and emits the pair: *<target*, list(*source*)*>*
  ◼ For a given web page, we know what pages link to it

# What about the following problems?

- ☐ How do we assign work units to workers?
- ☐ What if we have more work units than workers?
- ☐ What if workers need to share partial results?
- ☐ How do we aggregate partial results?
- ☐ How do we know all the workers have finished?
- ☐ What if workers die?

# MapReduce Runtime

- ☐ Handles scheduling
  - ■ Assigns workers to map and reduce tasks
- ☐ Handles "data distribution"
  - ■ Moves the process to the data
- ☐ Handles synchronization
  - ■ Gathers, sorts, and shuffles intermediate data
- ☐ Handles faults
  - ■ Detects worker failures and restarts
- ☐ Everything happens on top of a distributed file system

# How do we get data to the workers?

**NAS (Network Attached Storage)**

**Compute Nodes**

**What's the problem here?**

**SAN (Storage Area Network)**

# Distributed File System

- Don't move data to workers... Move workers to the data!
  - Store data on the local disks for nodes in the cluster
  - Start up the workers on the node that has the data local
- Why?
  - Not enough RAM to hold all the data in memory
  - Disk access is slow, disk throughput is good
- A distributed file system is the answer
  - GFS (Google File System)
  - HDFS for Hadoop

# Outline

- Introduction to MapReduce
  - Google File System
- The Hadoop Framework

# GFS: Assumptions

- Commodity hardware over "exotic" hardware
- High component failure rates
  - Inexpensive commodity components fail all the time
- "Modest" number of HUGE files
- Files are write-once, mostly appended to
- Perhaps concurrently
- Large streaming reads over random access
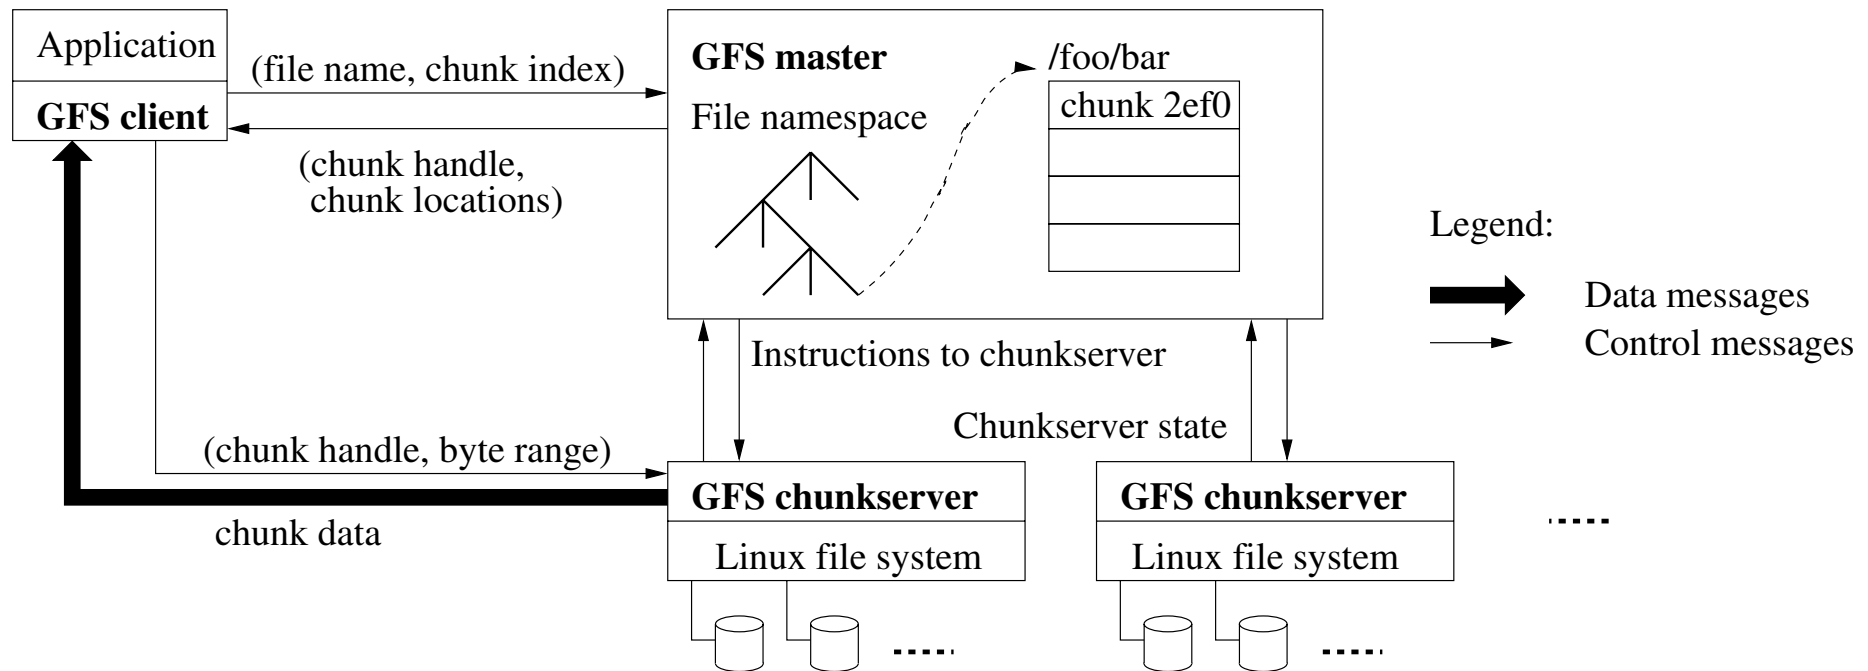- High sustained throughput over low latency

# GFS: Design Decisions

- ## Files stored as chunks
  - ### Fixed size (64MB)
- ## Reliability through replication
  - ### Each chunk replicated across 3+ chunkservers
- ## Single master to coordinate access, keep metadata
  - ### Simple centralized management
- ## No data caching
  - ### Little benefit due to large data sets, streaming reads
- ## Simplify the API
  - ### Push some of the issues onto the client (e.g., data layout)

# GFS Architecture

Application

(file name, chunk index)

**GFS client**

(chunk handle,
chunk locations)

**GFS master**

File namespace

/foo/bar

chunk 2ef0

Instructions to chunkserver

Chunkserver state

(chunk handle, byte range)

**GFS chunkserver**

Linux file system

**GFS chunkserver**

Linux file system

.....

chunk data

Legend:

Data messages

Control messages

# Single Master

- ## We know this is a:
    - ### Single point of failure
    - ### Scalability bottleneck
- ## GFS solutions:
    - ### Shadow masters
    - ### Minimize master involvement
        - Never move data through it, use only for metadata (and cache metadata at clients)
        - Large chunk size
        - Master delegates authority to primary replicas in data mutations (chunk leases)
- ## Simple, and good enough!

# Master's Responsibilities (1/2)

- Metadata storage

- Namespace management/locking

- Periodic communication with chunkservers
  - Give instructions, collect state, track cluster health

- Chunk creation, re-replication, rebalancing
  - Balance space utilization and access speed
  - Spread replicas across racks to reduce correlated failures
  - Re-replicate data if redundancy falls below threshold
  - Rebalance data to smooth out storage and request load

# Master's Responsibilities (2/2)

- ## Garbage Collection
  - Simpler, more reliable than traditional file delete
  - Master logs the deletion, renames the file to a hidden name
  - Lazily garbage collects hidden files
- ## Stale replica deletion
  - Detect "stale" replicas using chunk version numbers

# Metadata

- Global metadata is stored on the master
  - File and chunk namespaces
  - Mapping from files to chunks
  - Locations of each chunk's replicas
- All in memory (64 Mbytes / chunk)
  - Fast
  - Easily accessible
- Master has an operation log for persistent logging of critical metadata updates
  - Persistent on local disk
  - Replicated
  - Checkpoints for faster recovery

# Coordination

- ☐ Master data structures
    - ■ Task status: (idle, in-progress, completed)
    - ■ Idle tasks get scheduled as workers become available
    - ■ When a map task completes, it sends the master the location and sizes of its R intermediate files, one for each reducer
    - ■ Master pushes this info to reducers
- ☐ Master pings workers periodically to detect failures

# Failures

- ☐ Map worker failure
  - ■ Map tasks completed or in-progress at worker are reset to idle
  - ■ Reduce workers are notified when task is rescheduled on another worker
- ☐ Reduce worker failure
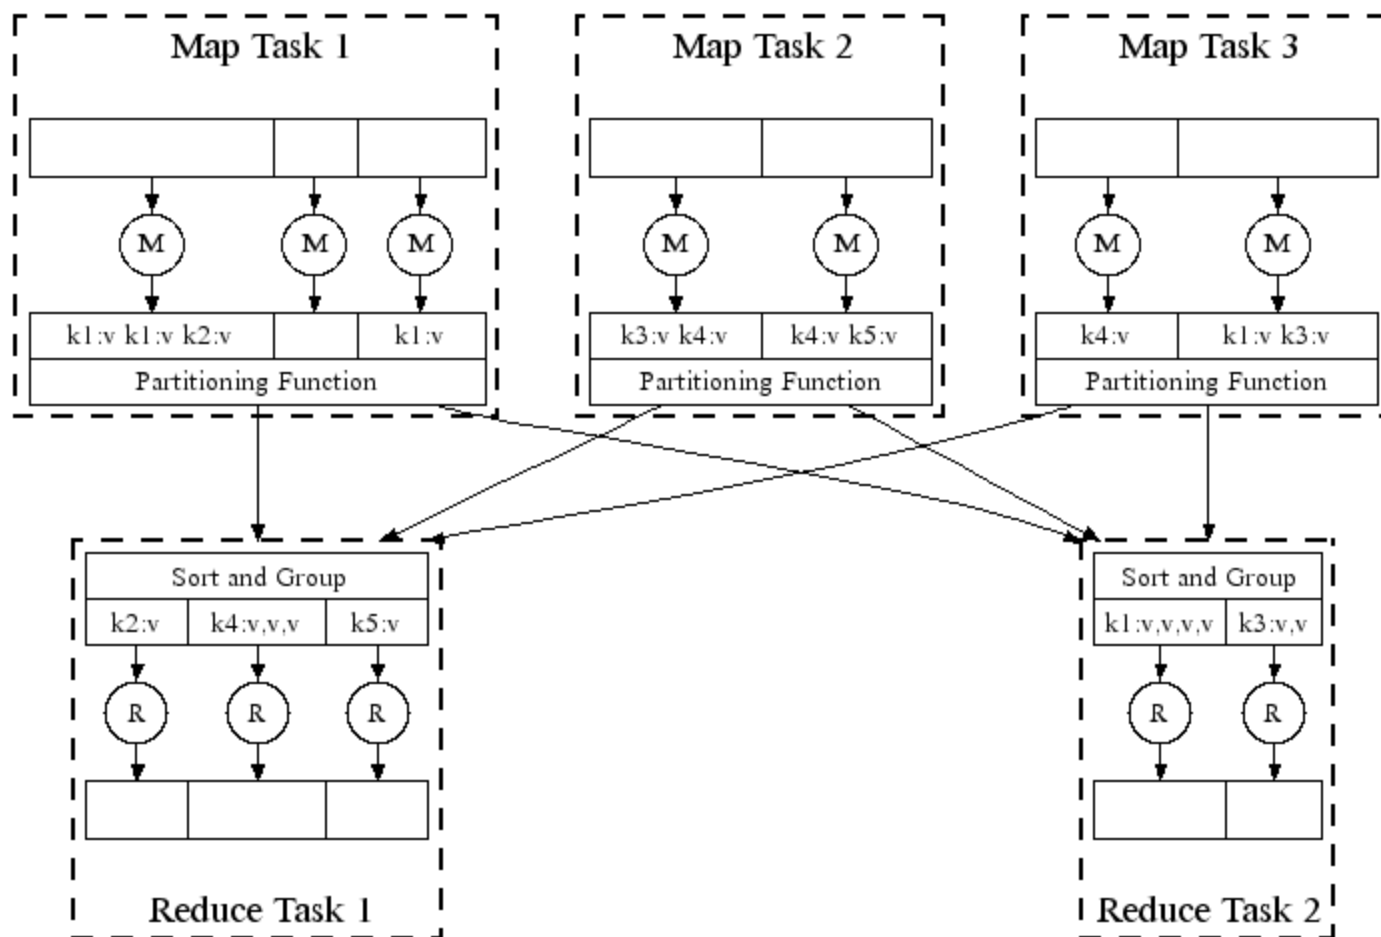  - ■ Only in-progress tasks are reset to idle
- ☐ Master failure
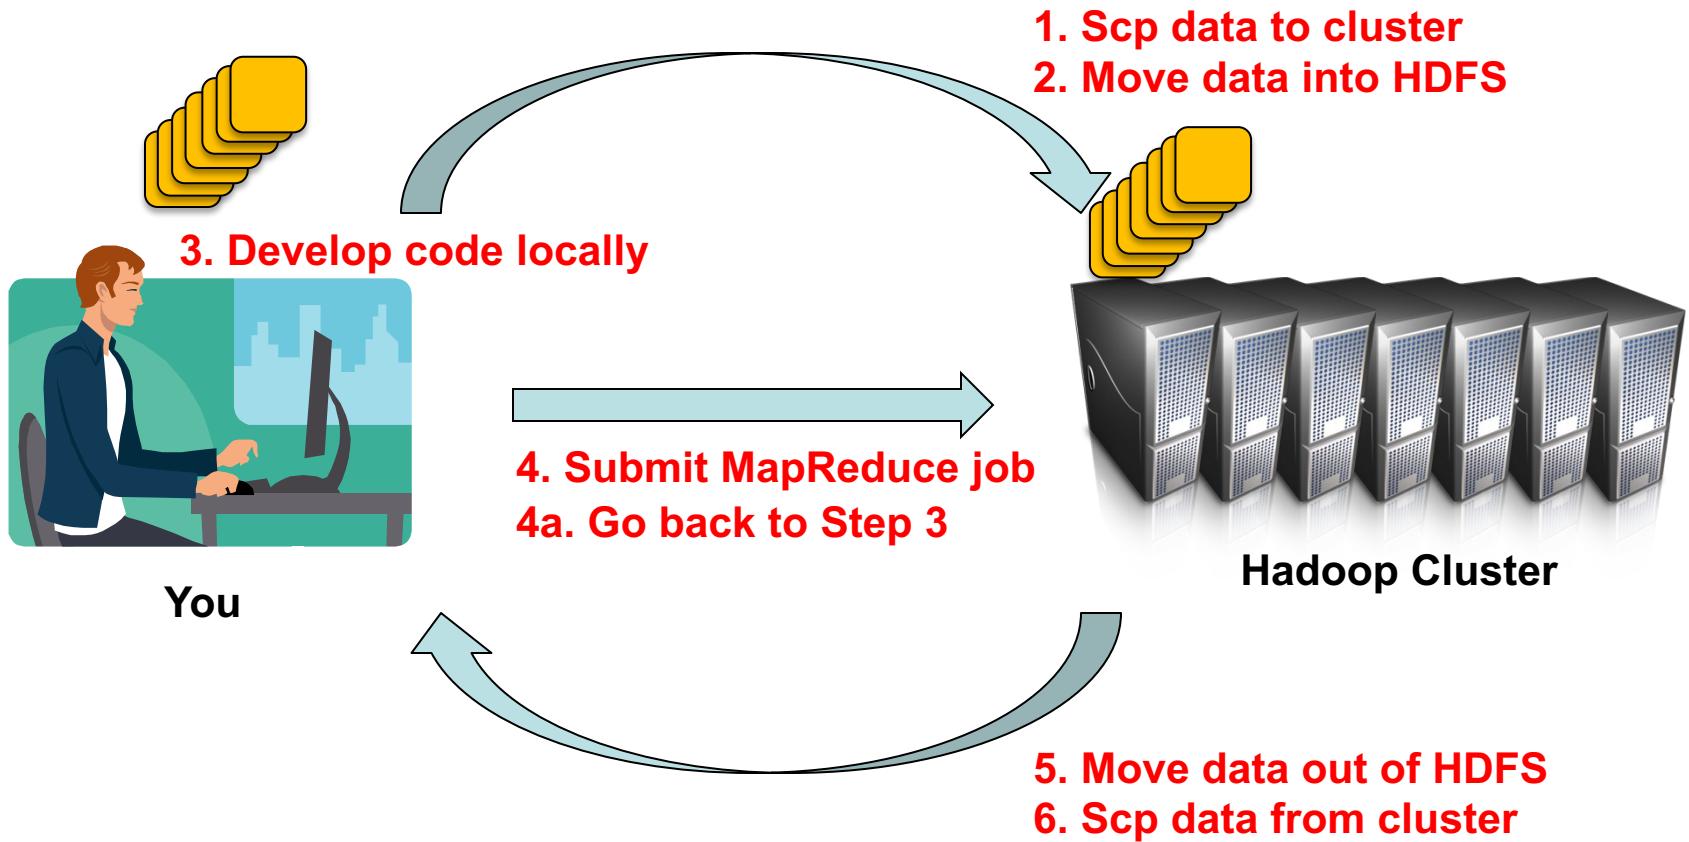  - ■ MapReduce task is aborted and client is notified

# Execution
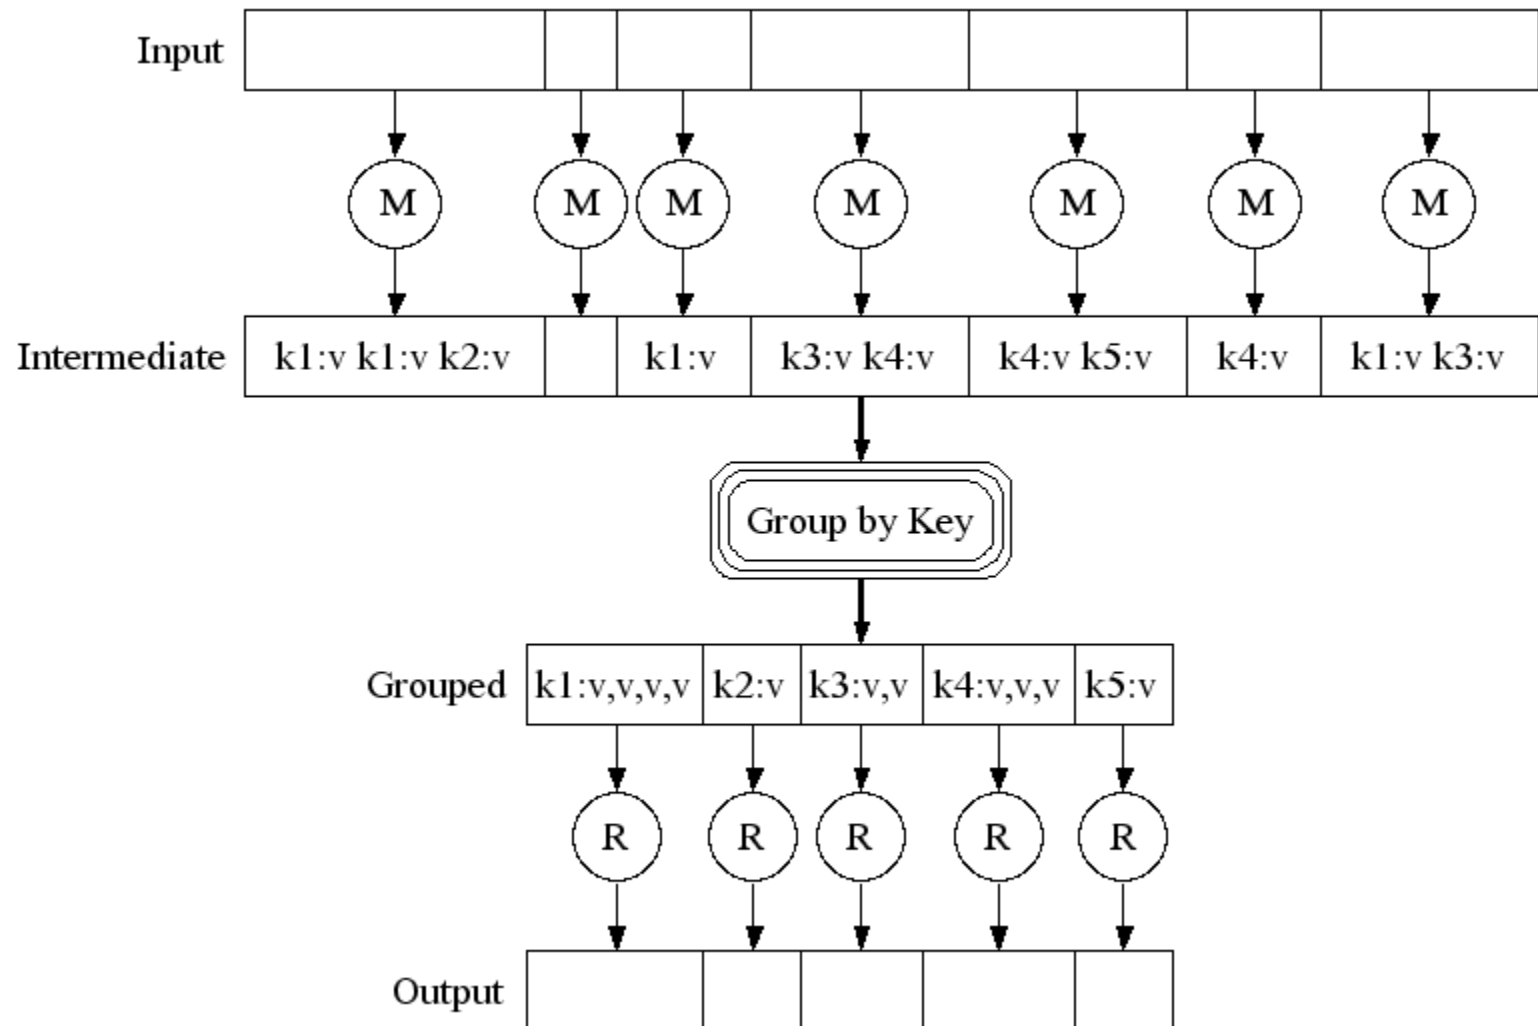
# Parallel Execution

# How many Map and Reduce jobs?

- ☐ M map tasks, R reduce tasks
- ☐ Rule of thumb:
    - ■ Make M and R much larger than the number of nodes in cluster
    - ■ One DFS chunk per map is common
    - ■ Improves dynamic load balancing and speeds recovery from worker failure
- ☐ Usually R is smaller than M, because output is spread across R files

# From Theory to Practice



**1. Scp data to cluster**
**2. Move data into HDFS**

**3. Develop code locally**

**4. Submit MapReduce job**
**4a. Go back to Step 3**

**You**

**Hadoop Cluster**

**5. Move data out of HDFS**
**6. Scp data from cluster**

# Execution Summary

☐ How is this distributed?

1. Partition input key/value pairs into chunks, run map() tasks in parallel

2. After all map()s are complete, consolidate all emitted values for each unique emitted key

3. Now partition space of output map keys, and run reduce() in parallel

☐ If map() or reduce() fails, reexecute!

# Outline

- Introduction to MapReduce
  - Google File System
- The Hadoop Framework

# Hadoop

- Hadoop (High-availability distributed object-oriented platform)

- An open-source software framework that supports data-intensive distributed applications

- Distributed File System (HDFS)

  - Runs on same machines

  - Replicates data 3x for fault-tolerance

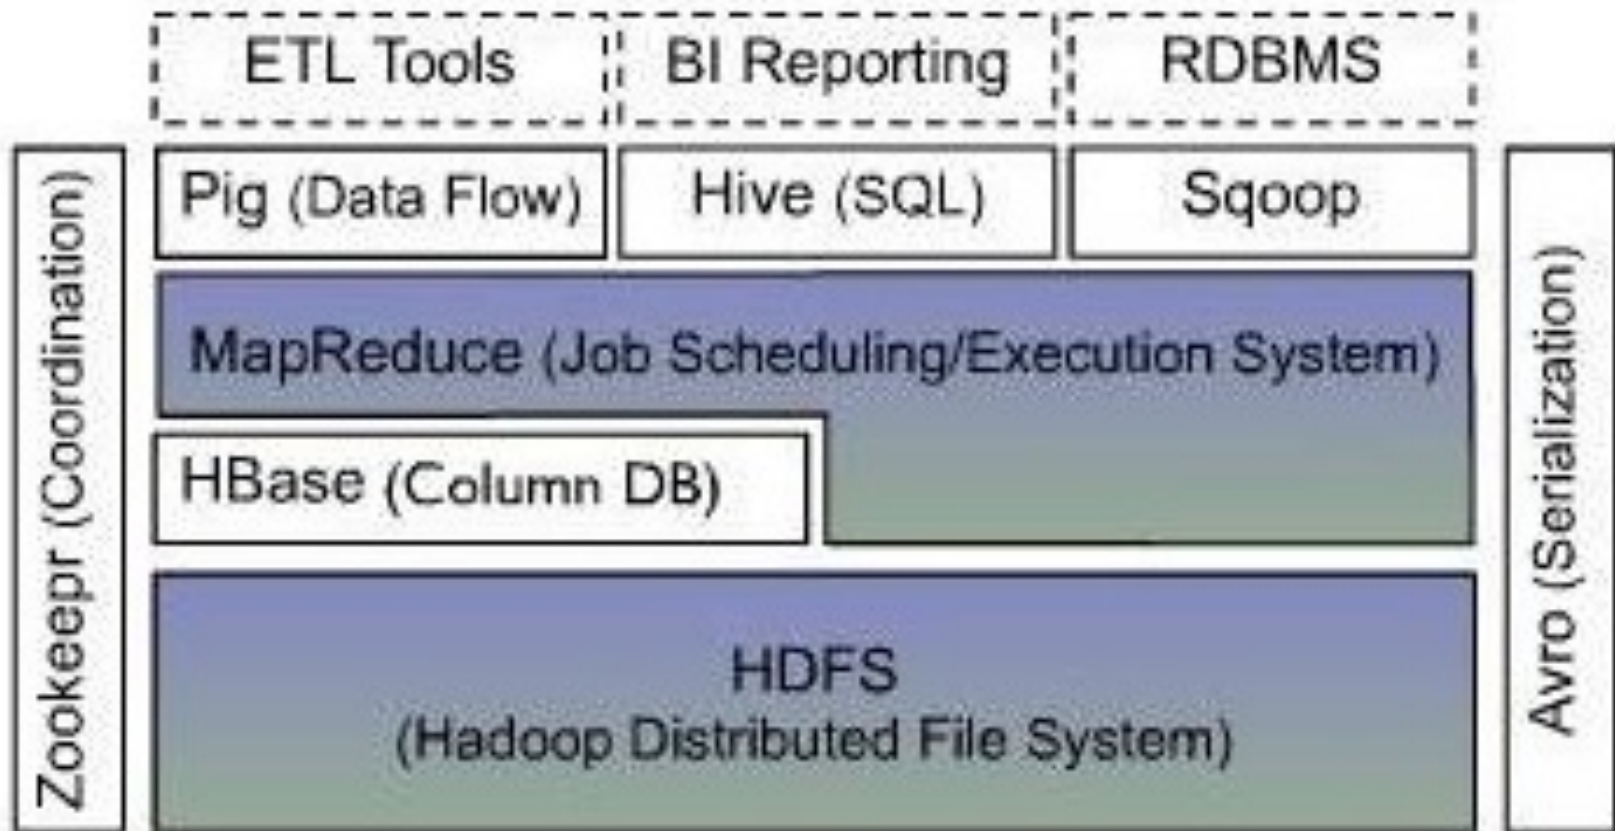- It supports the running of applications on large clusters of commodity hardware

# Hadoop

- Hadoop was derived from Google's MapReduce and Google File System (GFS) papers

- Hadoop is written in the Java programming language

  - Download from: http://hadoop.apache.org/

# Hadoop Ecosystem



Source: Cloudera

# Terminology

| Google calls it: | Hadoop equivalent: |
|---|---|
| MapReduce | Hadoop |
| GFS | HDFS |
| Bigtable | HBase |
| Chubby | Zookeeper |

# Some MapReduce Terminology

- *Job* – A "full program" - an execution of a Mapper and Reducer across a data set

- *Task* – An execution of a Mapper or a Reducer on a slice of data
  - a.k.a. Task-In-Progress (TIP)

- Task Attempt – A particular instance of an attempt to execute a task on a machine

# Task Attempts

- A particular task will be attempted at least once, possibly more times if it crashes

  - If the same input causes crashes over and over, that input will eventually be abandoned

- Multiple attempts at one task may occur in parallel with speculative execution turned on

  - Task ID from *TaskInProgress* is not a unique identifier; don't use it that way

# MapReduce: High Level



**In our case:** pp1.cs.nctu.edu.tw
**or hadoop.nchc.org.tw**
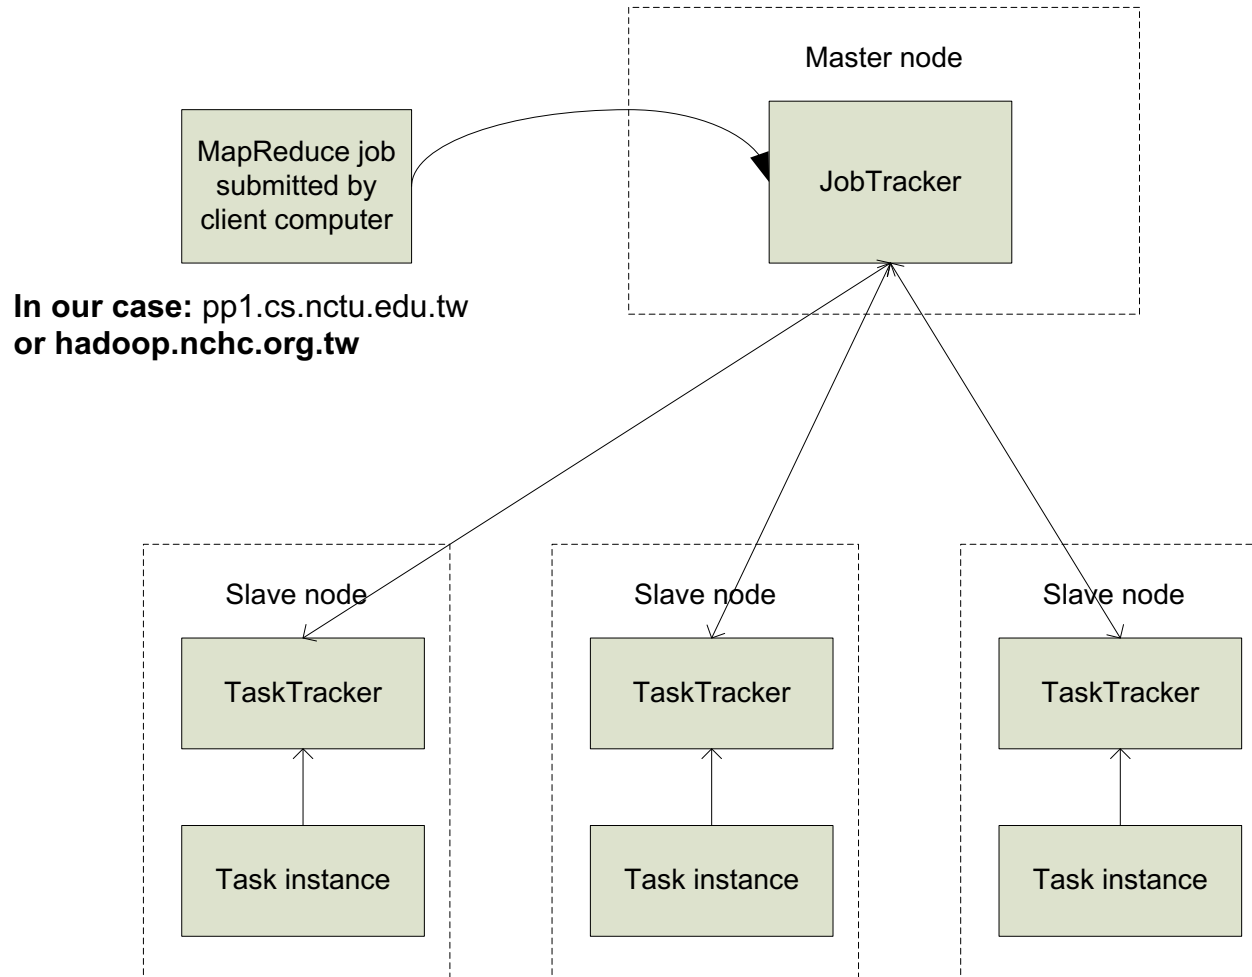
# Nodes, Trackers, Tasks

- Master node runs *JobTracker* instance, which accepts *Job* requests from clients

- *TaskTracker* instances run on slave nodes

- TaskTracker forks separate Java processes for task instances

# Job Distribution

- MapReduce programs are contained in a Java "jar" file + an XML file containing serialized program configuration options

- Running a MapReduce job places these files into the HDFS and notifies TaskTrackers where to retrieve the relevant program code

- … Where's the data distribution?

# Data Distribution

- ## Implicit in design of MapReduce!
  - All mappers are equivalent; so map whatever data is local to a particular node in HDFS
- ## If lots of data does happen to pile up on the same node, nearby nodes will map instead
  - Data transfer is handled implicitly by HDFS

# What Happens In Hadoop? Depth First

# Job Launch Process: Client

- **Client program creates a *JobConf***
  - Identify classes implementing *Mapper* and *Reducer* interfaces
    - JobConf.setMapperClass(), setReducerClass()
  - Specify inputs, outputs
    - FileInputFormat.setInputPath(),
    - FileOutputFormat.setOutputPath()
  - Optionally, other options too:
    - JobConf.setNumReduceTasks(), JobConf.setOutputFormat()…

# Job Launch Process: *JobClient*

- Pass JobConf to JobClient.runJob() or submitJob()
  - runJob() blocks, submitJob() does not
- *JobClient*:
  - Determines proper division of input into *InputSplits*
  - Sends job data to master *JobTracker* server

# Job Launch Process: *JobTracker*

- *JobTracker*:
  - Inserts jar and JobConf (serialized to XML) in shared location
  - Posts a *JobInProgress* to its run queue

# Job Launch Process: *TaskTracker*

- *TaskTrackers* running on slave nodes periodically query *JobTracker* for work

- Retrieve job-specific jar and config

- Launch task in separate instance of Java
  - main() is provided by Hadoop

# Job Launch Process: Task

- TaskTracker.Child.main():
  - ❑ Sets up the child *TaskInProgress* attempt
  - ❑ Reads XML configuration
  - ❑ Connects back to necessary MapReduce components via RPC
  - ❑ Uses *TaskRunner* to launch user process

# Job Launch Process: *TaskRunner*

- *TaskRunner*, *MapTaskRunner*, *MapRunner* work in a daisy-chain to launch your *Mapper*
  - Task knows ahead of time which *InputSplits* it should be mapping
  - Calls *Mapper* once for each record retrieved from the InputSplit
- Running the *Reducer* is much the same

# Creating the *Mapper*

- You provide the instance of *Mapper*
  - Should extend *MapReduceBase*
- One instance of your Mapper is initialized by the *MapTaskRunner* for a *TaskInProgress*
  - Exists in separate process from all other instances of Mapper – no data sharing!

# *Mapper*

- void map(K1 key,

    V1 value,

    OutputCollector<K2, V2> output,

    Reporter reporter)


- *K* types implement *WritableComparable*
- *V* types implement *Writable*

# What is Writable?

- Hadoop defines its own "box" classes for strings *(Text),* integers *(IntWritable)*, etc.

- All values are instances of *Writable*

- All keys are instances of *WritableComparable*

# Getting Data To The Mapper
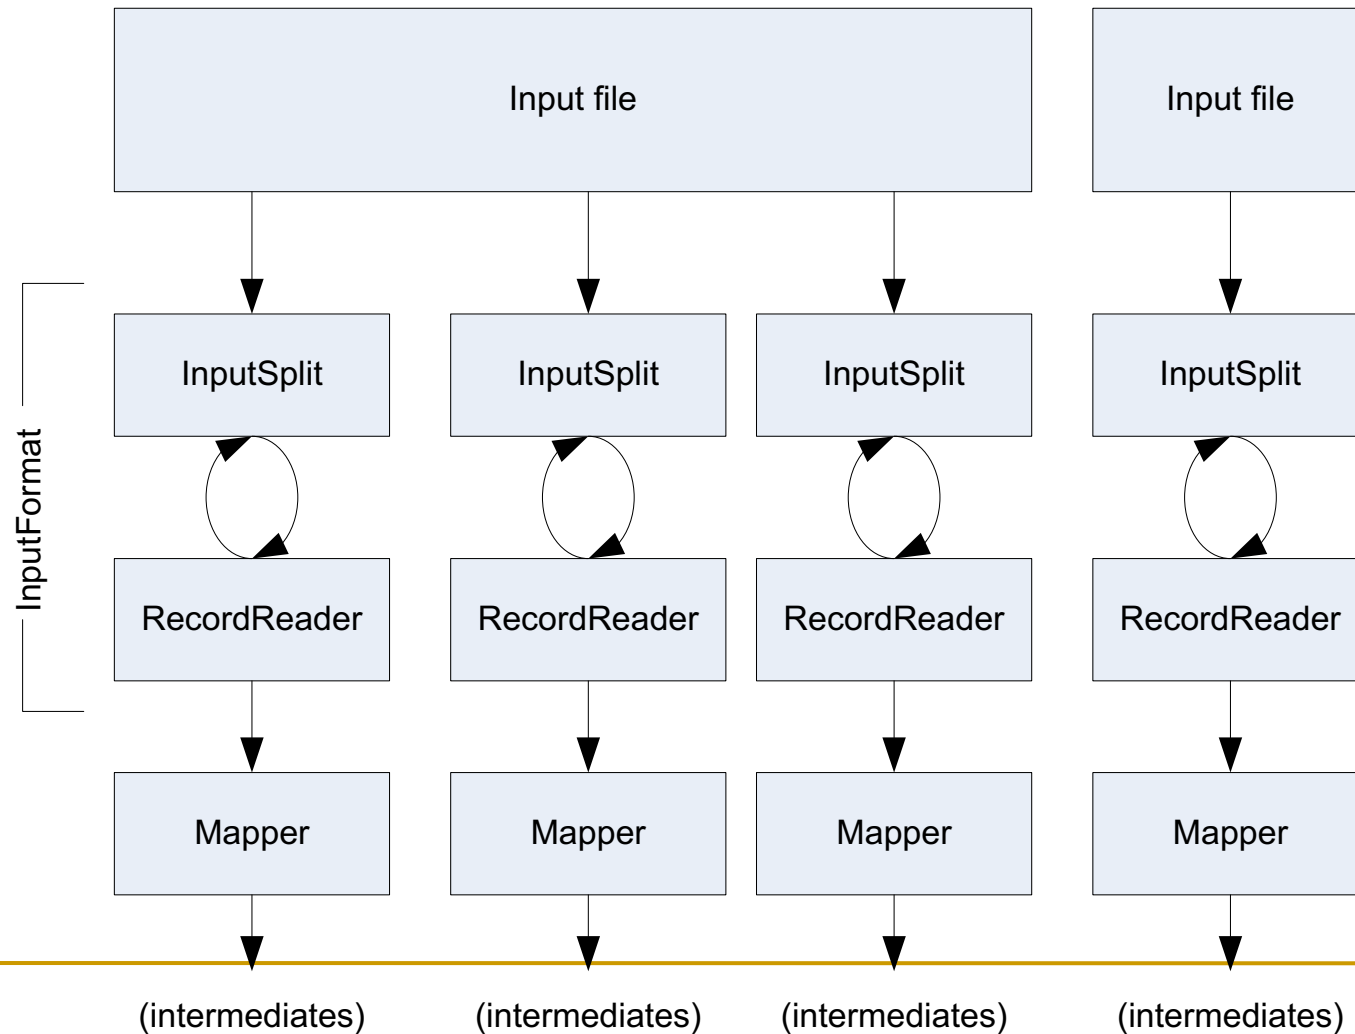
# Reading Data

- Data sets are specified by *InputFormats*
  - Defines input data (e.g., a directory)
  - Identifies partitions of the data that form an *InputSplit*
  - Factory for *RecordReader* objects to extract (k, v) records from the input source

# *FileInputFormat* and Friends

- *TextInputFormat* – Treats each '\n'-terminated line of a file as a value

- *KeyValueTextInputFormat* – Maps '\n'-terminated text lines of "k SEP v"

- *SequenceFileInputFormat* – Binary file of (k, v) pairs with some add'l metadata

- *SequenceFileAsTextInputFormat* – Same, but maps (k.toString(), v.toString())

# Filtering File Inputs

- *FileInputFormat* will read all files out of a specified directory and send them to the mapper

- Delegates filtering this file list to a method subclasses may override

  - *e.g.,* Create your own "xyzFileInputFormat" to read *.xyz from directory list

# Record Readers

- Each *InputFormat* provides its own *RecordReader* implementation
  - Provides (unused?) capability multiplexing
- *LineRecordReader* – Reads a line from a text file
- *KeyValueRecordReader* – Used by *KeyValueTextInputFormat*

# Input Split Size

- *FileInputFormat* will divide large files into chunks

  - Exact size controlled by mapred.min.split.size

- RecordReaders receive file, offset, and length of chunk

- Custom *InputFormat* implementations may override split size – e.g., "NeverChunkFile"

# Sending Data To Reducers

- Map function receives *OutputCollector* object
  - OutputCollector.collect() takes (k, v) elements
- Any *(WritableComparable, Writable)* can be used
- By default, mapper output type assumed to be same as reducer output type

# *WritableComparator*

- Compares WritableComparable data
  - ❑ Will call WritableComparable.compare()
  - ❑ Can provide fast path for serialized data
- *JobConf.setOutputValueGroupingComparator()*

# Sending Data To The Client

- *Reporter* object sent to Mapper allows simple asynchronous feedback
  - incrCounter(Enum key, long amount)
  - setStatus(String msg)
- Allows self-identification of input
  - InputSplit getInputSplit()

# Partition And Shuffle
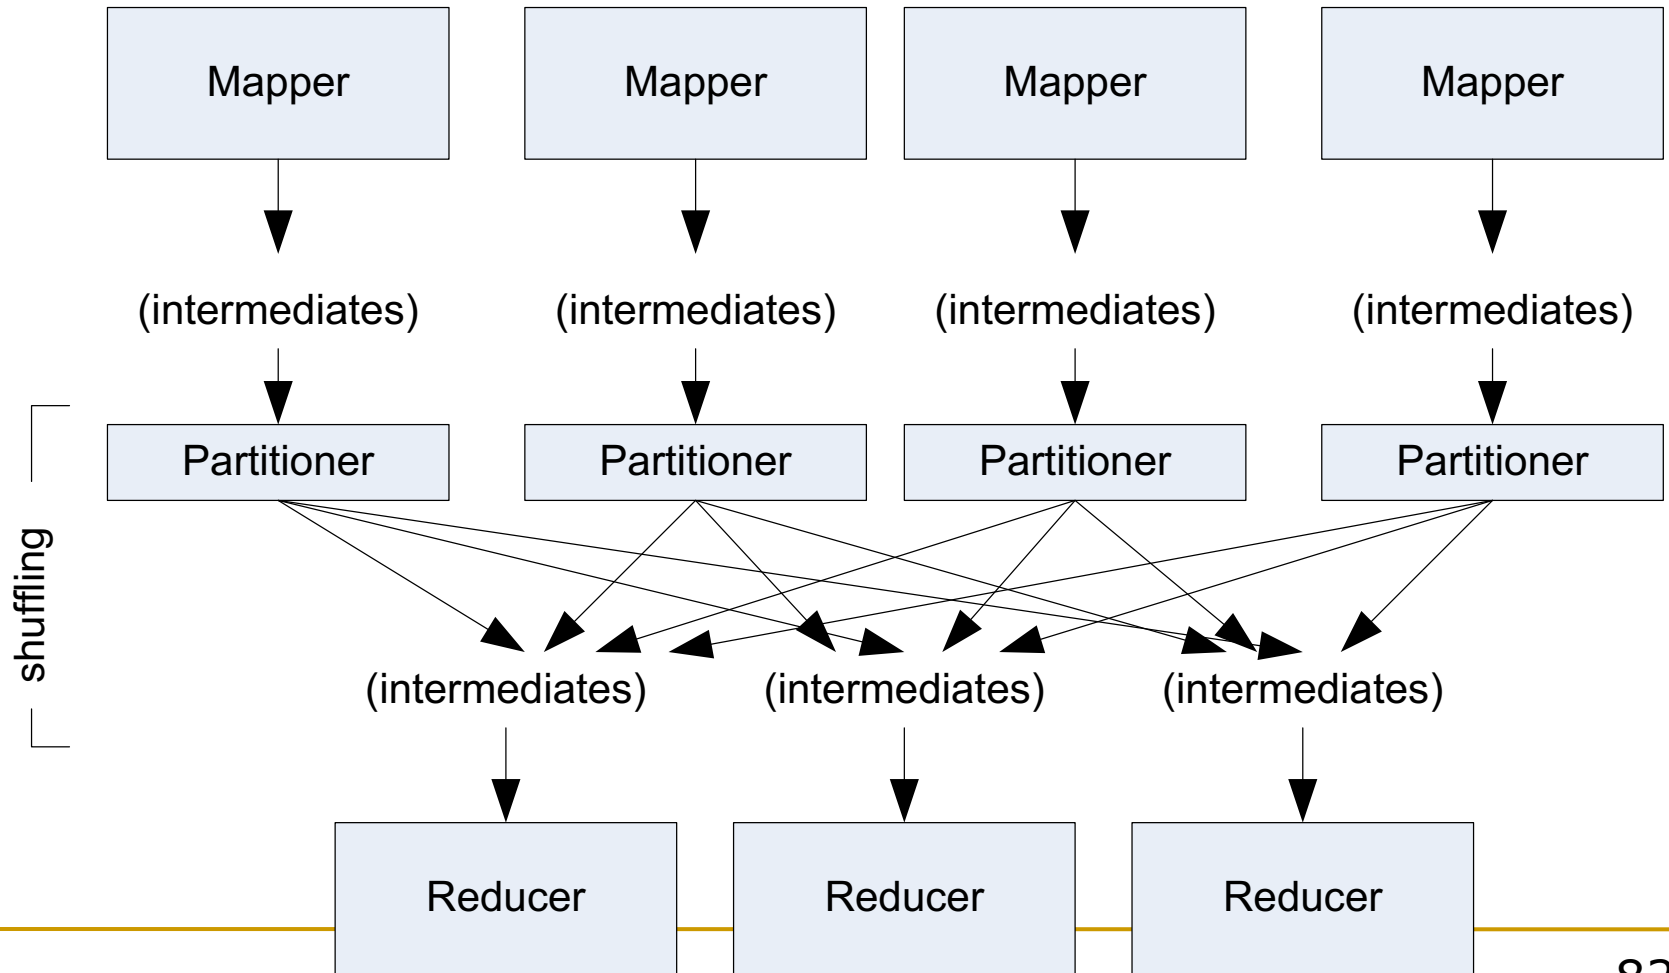
# *Partitioner*

- int getPartition(key, val, numPartitions)
  - ❑ Outputs the partition number for a given key
  - ❑ One partition == values sent to one Reduce task
- *HashPartitioner* used by default
  - ❑ Uses key.hashCode() to return partition num
- *JobConf* sets *Partitioner* implementation

# Reduction

- reduce(  K2 key,

    Iterator<V2> values,

    OutputCollector<K3, V3> output,

    Reporter reporter )

- Keys & values sent to one partition all go to the same reduce task

- Calls are sorted by key – "earlier" keys are reduced and output before "later" keys

# Finally: Writing The Output
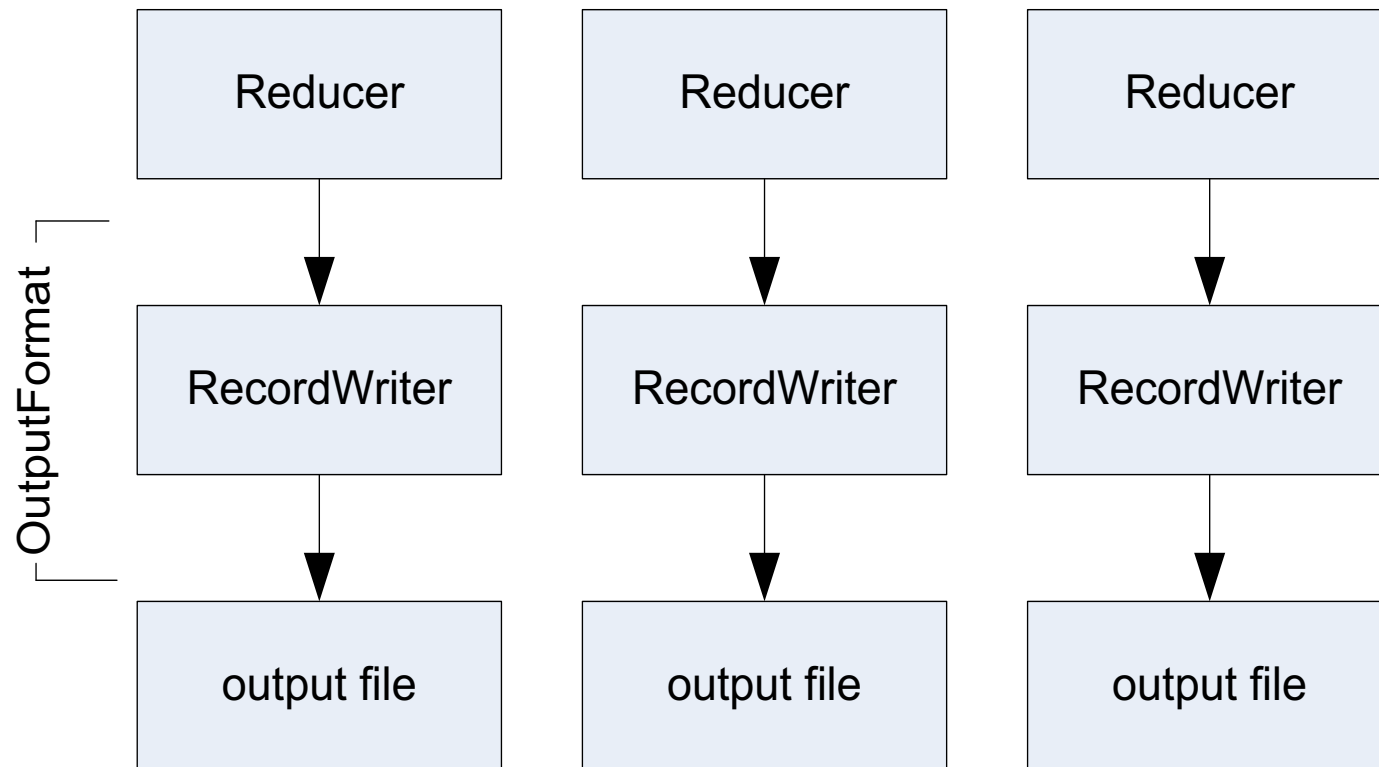
# *OutputFormat*

- Analogous to *InputFormat*

- *TextOutputFormat* – Writes "key val\n" strings to output file

- *SequenceFileOutputFormat* – Uses a binary format to pack (k, v) pairs

- *NullOutputFormat* – Discards output
  - Only useful if defining own output methods within reduce()

# Methods to write MapReduce Jobs

- Typical – usually written in Java
  - MapReduce 2.0 API
  - MapReduce 1.0 API
- Pipes
  - Often used with C++
- Streaming
  - Uses stdin and stdout
  - Can use any language to write Map and Reduce functions
- Abstraction libraries
  - Hive, Pig, etc.
  - Written in a higher level language, generate one or more MapReduce jobs

# Example Program - Wordcount

- **map()**
  - Receives a chunk of text
  - Outputs a set of word/count pairs
- **reduce()**
  - Receives a key and all its associated values
  - Outputs the key and the sum of the values

```
package org.myorg;
import java.io.IOException;
import java.util.*;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.conf.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.util.*;

public class WordCount {
```

# Wordcount – main( )

```
public static void main(String[] args) throws Exception {
    JobConf conf = new JobConf(WordCount.class);
    conf.setJobName("wordcount");
    conf.setOutputKeyClass(Text.class);
    conf.setOutputValueClass(IntWritable.class);

    conf.setMapperClass(Map.class);
    conf.setReducerClass(Reduce.class);

    conf.setInputFormat(TextInputFormat.class);
    conf.setOutputFormat(TextOutputFormat.class);

    FileInputFormat.setInputPaths(conf, new Path(args[0]));
    FileOutputFormat.setOutputPath(conf, new Path(args[1]));

    JobClient.runJob(conf);
}
```

# Wordcount – map( )

```
public static class Map extends MapReduceBase … {
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(LongWritable key, Text value,
                    OutputCollector<Text, IntWritable> output, …) … {
        String line = value.toString();
        StringTokenizer tokenizer = new StringTokenizer(line);
        while (tokenizer.hasMoreTokens()) {
            word.set(tokenizer.nextToken());
            output.collect(word, one);
        }
    }
}
```

# Wordcount – reduce( )

```
public static class Reduce extends MapReduceBase … {
    public void reduce(Text key, Iterator<IntWritable> values,
                        OutputCollector<Text, IntWritable> output, …) … {
        int sum = 0;
        while (values.hasNext()) {
            sum += values.next().get();
        }
        output.collect(key, new IntWritable(sum));
    }
}
}
```

# Hadoop Pipes: wordcount – main()

- **Allows you to create map/reduce jobs in C++**

```
#include "hadoop/Pipes.hh"
#include "hadoop/TemplateFactory.hh"
#include "hadoop/StringUtils.hh"

int main(int argc, char *argv[]) {
  return
HadoopPipes::runTask(HadoopPipes::TemplateFactory<WordCountMap,
                WordCountReduce>());
}
```

# Hadoop Pipes: wordcount – map()

- **Allows you to create map/reduce jobs in C++**

```cpp
class WordCountMap: public HadoopPipes::Mapper {
public:
  WordCountMap(HadoopPipes::TaskContext& context){}
  void map(HadoopPipes::MapContext& context) {
    std::vector<std::string> words =
      HadoopUtils::splitString(context.getInputValue(), " ");
    for(unsigned int i=0; i < words.size(); ++i) {
      context.emit(words[i], "1");
    }
  }
};
```

# Hadoop Pipes: wordcount – reduce()

- **Allows you to create map/reduce jobs in C++**

```cpp
class WordCountReduce: public HadoopPipes::Reducer {
public:
  WordCountReduce(HadoopPipes::TaskContext& context){}
  void reduce(HadoopPipes::ReduceContext& context) {
    int sum = 0;
    while (context.nextValue()) {
      sum += HadoopUtils::toInt(context.getInputValue());
    }
    context.emit(context.getInputKey(), HadoopUtils::toString(sum));
  }
};
```

# Running with Hadoop Pipes

- Allows you to create map/reduce jobs in C++

g++ wordcount.cpp -o wordcount -lhadooputils \
-lhadooppipes -lpthread


hadoop pipes -D hadoop.pipes.java.recordreader=true \
-D hadoop.pipes.java.recordwriter=true \
-program examples/bin/wordcount \
-input examples/input -output examples/output

# Hadoop Streaming

- Allows you to create and run map/reduce jobs with any executable

- Similar to unix pipes, e.g.:
  - format is: Input | Mapper | Reducer
  - echo "this sentence has five lines" | cat | wc -w

```
$HADOOP_HOME/bin/hadoop jar $HADOOP_HOME/hadoop-streaming.jar \
        -input myInputDirs \
        -output myOutputDir \
        -mapper /bin/cat \
        -reducer /bin/wc
```

# Hadoop Streaming

- Mapper and Reducer receive data from stdin and output to stdout

- Hadoop takes care of the transmission of data between the map/reduce tasks
  - It is still the programmer's responsibility to set the correct key/value
  - Default format: "key \t value\n"

- Let's look at a C++/Python example of a MapReduce word count program…

# wordcount_map.cpp

```cpp
#include<iostream>
using namespace std;

int main()
{
    string word;
    while (cin >> word)
    {
        cout << word << "\t" << "1" << endl;
    }
    return 0;
}
```

# Streaming_Mapper.py

```python
# read in one line of input at a time from stdin
for line in sys.stdin:
    line = line.strip()        # string
    words = line.split()       # list of strings

    # write data on stdout
    for word in words:
        print '%s\t%i' % (word, 1)
```

# Hadoop Streaming

- **What are we outputting?**
  - Example output: "the        1"
  - By default, "the" is the key, and "1" is the value
- **Hadoop Streaming handles delivering this key/value pair to a Reducer**
  - Able to send similar keys to the same Reducer or to an intermediary Combiner

# wordcount_reduce.cpp

```cpp
#include<vector>
#include<map>
#include<string>
#include<iostream>
using namespace std;
int main() {
    string key;
    string value;
    map<string,int> mapTemp;
    while(cin >> key >> value) {
        if(mapTemp.find(key) !=
    mapTemp.end())
            mapTemp[key] += 1;
        else
            mapTemp[key] = 1;
    }

    map<string,int>::iterator it = mapTemp.begin();
    for(; it != mapTemp.end(); ++it) {
        cout << it->first << "\t" << it->second << endl;
    }


    return 0;
}
```

# Streaming_Reducer.py

```
wordcount = { }            # empty dictionary
# read in one line of input at a time from stdin
for line in sys.stdin:
    line = line.strip()            # string
    key,value = line.split()
    wordcount[key] = wordcount.get(key, 0) + value

    # write data on stdout
    for word, count in sorted(wordcount.items()):
        print '%s\t%i' % (word, count)
```

# Hadoop Streaming Gotcha

- **Streaming Reducer receives single lines (which are key/value pairs) from stdin**

  - Regular Reducer receives a collection of all the values for a particular key

  - It is still the case that all the values for a particular key will go to a single Reducer

# Using Hadoop Distributed File System (HDFS)

- Can access HDFS through various shell commands (see Further Resources slide for link to documentation)

  - hadoop fs –put <localsrc> … <dst>

  - hadoop fs –get <src> <localdst>

  - hadoop fs –ls

  - hadoop fs –rm file

# Configuring Number of Tasks

- **Normal method**
  - ❏ jobConf.setNumMapTasks(400)
  - ❏ jobConf.setNumReduceTasks(4)
- **Hadoop Streaming method**
  - ❏ -jobconf mapred.map.tasks=400
  - ❏ -jobconf mapred.reduce.tasks=4
- **Note: # of map tasks is only a hint to the framework. Actual number depends on the number of InputSplits generated**

# Running a Hadoop Job

- **Place input file into HDFS:**
  - hadoop fs –put ./input-file input-file
- **Run either normal or streaming version:**
  - hadoop jar Wordcount.jar org.myorg.Wordcount input-file output-file
  - hadoop jar hadoop-streaming.jar \
    -input input-file \
    -output output-file \
    -file Streaming_Mapper.py \
    -mapper python Streaming_Mapper.py \
    -file Streaming_Reducer.py \
    -reducer python Streaming_Reducer.py \

# Output Parsing

- Output of the reduce tasks must be retrieved:
  - hadoop fs –get output-file hadoop-output
- This creates a directory of output files, 1 per reduce task
  - Output files numbered part-00000, part-00001, etc.
- Sample output of Wordcount
  - head –n5 part-00000

        tis        1
        come       2
        coming     1
        edwin      1
        found      1

# Extra Output

- The stdout/stderr streams of Hadoop itself will be stored in an output file (whichever one is named in the startup script)
  - #$ -o output.$job_id

STARTUP_MSG: Starting NameNode
STARTUP_MSG:   host = svc-3024-8-10.rc.usf.edu/10.250.4.205
…
11/03/02 18:28:47 INFO mapred.FileInputFormat: Total input paths to process : 1
11/03/02 18:28:47 INFO mapred.JobClient: Running job: job_local_0001
…
11/03/02 18:28:48 INFO mapred.MapTask: numReduceTasks: 1
…
11/03/02 18:28:48 INFO mapred.TaskRunner: Task 'attempt_local_0001_m_000000_0' done.
11/03/02 18:28:48 INFO mapred.Merger: Merging 1 sorted segments
11/03/02 18:28:48 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 43927 bytes
11/03/02 18:28:48 INFO mapred.JobClient:  map 100% reduce 0%
…
11/03/02 18:28:49 INFO mapred.TaskRunner: Task 'attempt_local_0001_r_000000_0' done.
11/03/02 18:28:49 INFO mapred.JobClient: Job complete: job_local_0001

# Further Resources

- Hadoop Tutorial:
  http://developer.yahoo.com/hadoop/tutorial/module1.html
- Hadoop Streaming:
  http://hadoop.apache.org/common/docs/r0.15.2/streaming.html
- Hadoop API: http://hadoop.apache.org/common/docs/current/api
- HDFS Commands Reference:
  http://hadoop.apache.org/hdfs/docs/current/file_system_shell.html

# Public Hadoop Cluster

- ## http://hadoop.nchc.org.tw/

# Reference

- Running Hadoop on Ubuntu Linux (Single-Node Cluster)

  - http://www.michael-noll.com/tutorials/running-hadoop-on-ubuntu-linux-single-node-cluster/

# Conclusions

- MapReduce proven to be useful abstraction

- Greatly simplifies large-scale computations

- Fun to use:

  - Focus on problem

  - Let library deal with messy details